

## Processes

- One of the most important concepts in all modern operating systems
- A process is a **container** (or execution environment) for a program (software) in execution.
- Any software is executed with one or more processes.
  - Java VM (JVM), MS Word, Excel, PPT, Firefox, iTunes, Google Chrome, etc.

1

2

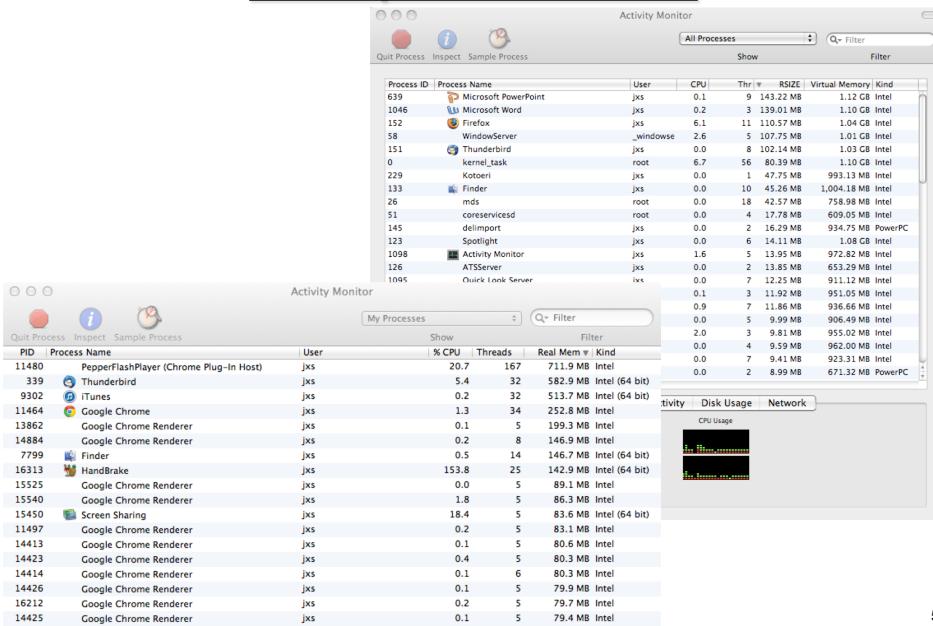
## Multi-Tasking (Time Sharing)

- All modern OSes support **multi-tasking**.
  - Doing **multiple things** with **multiple processes** on a **single CPU** (more precisely, a single CPU core)
    - e.g., writing a document with MS Word while playing a music with iTunes and downloading files with a Web browser
- At any moment, a single CPU core can execute a single program (process).
- **Pseudo concurrency/parallelism**: An **illusion** for human eyes as if multiple processes are executed at the same time.
  - OS periodically assigns one process to another to the CPU core.
    - every several tens of milliseconds to several hundreds of milliseconds
    - A CPU core is multiplexed among processes; processes are not executed in a completely parallel manner.

3

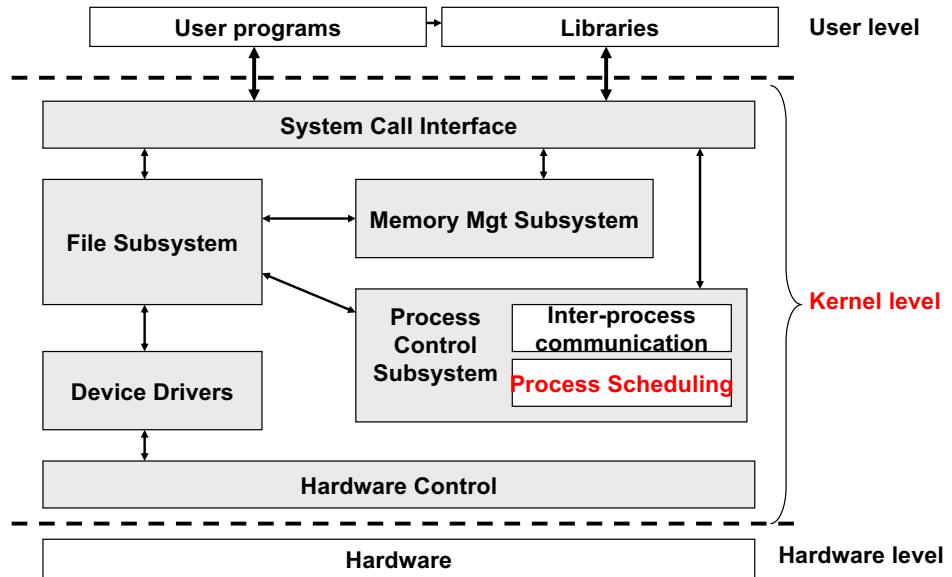
4

## Example Processes



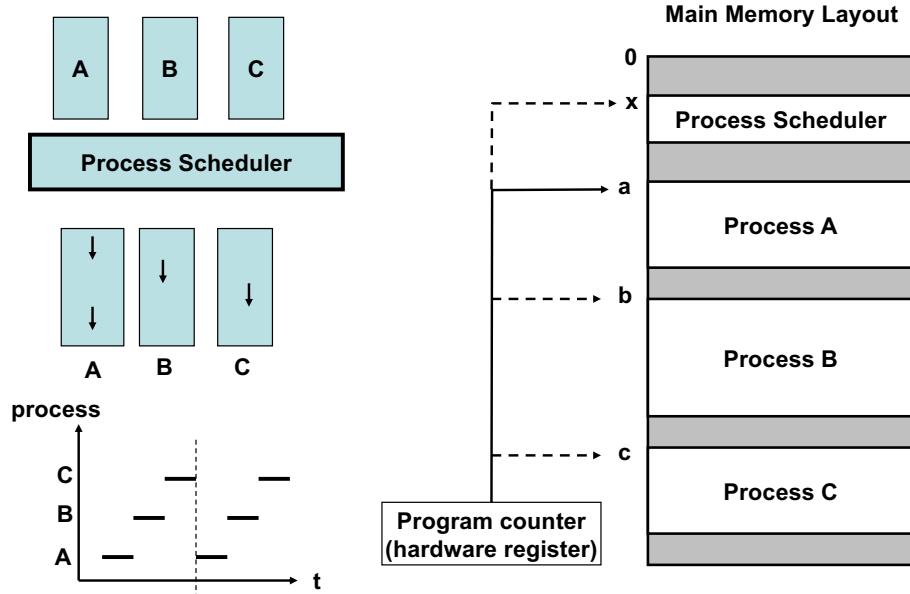
5

## An Architectural View of an OS



6

## Process Scheduling



7

## Why Threads?

- **Good old days:** Process-based, **coarse-grained** concurrency was enough
  - e.g., Editing a document while downloading files
- **Now:** **Fine-grained** concurrency is required.
  - Each program is required to execute multiple **in-program tasks**.
    - Doing different things at the same time.
    - Doing the same or a similar thing for multiple clients at the same time.

8

- Each program is required to do different things at the same time.

- Word
  - Displaying text and images
  - Responding to keystrokes and mouse inputs from the user
  - Downloading fancy document templates from MS web site
  - Performing spelling and grammar checking in the background

- Web browser
  - Displaying text and images
  - Responding to keystrokes/mouse inputs
  - Checking and downloading software updates
  - Multiple tabs

- iTunes
  - Playing music
  - Downloading music and its metadata (album's cover img, song titles, lyrics...)
  - Capturing music data (as .mp3 files, for example) from a CD

- # of threads
  - Kernel tasks: 121
  - Dropbox: 103
  - Firefox: 74
  - Thunderbird: 57
  - iTunes: 30
  - MS PPT: 21

9

10

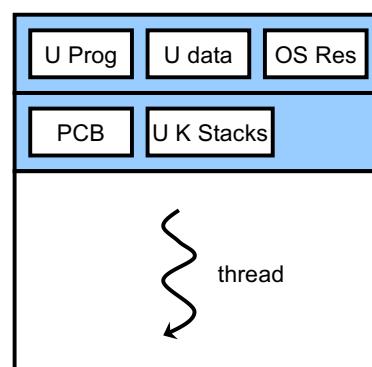
- Each program is required to do the same or a similar thing for multiple clients at the same time.

- Web server
  - Accepts and parses an HTTP request
  - Finds a target file
  - Makes an HTTP message (header, payload, etc.)
  - Returns a target file with the HTTP message

## • Assign **threads** to **in-program tasks**

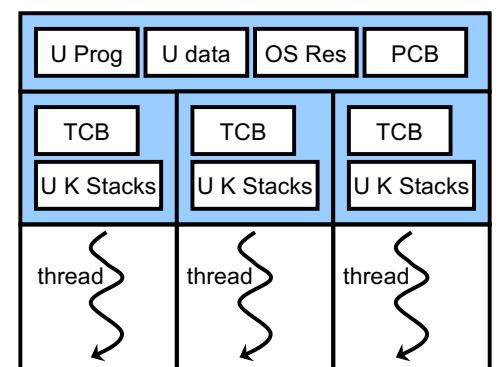
- Process-creation is heavyweight.

- Time-consuming and resource intensive.
  - Creating a process is 30 times slower than creating a thread.
  - Process switching is 5 times slower than thread switching.



Single-threaded (traditional) process

11



Multi-threaded process

12

## Summary: Why Threads?

- **Fine-grained pseudo parallelism:** An **illusion** for human eyes as if multiple threads are executed on a process at the same time.
  - OS periodically assigns one thread to another to the CPU core.
  - A CPU core is multiplexed among processes and threads; processes/threads are not executed in a completely parallel manner on a single CPU core.

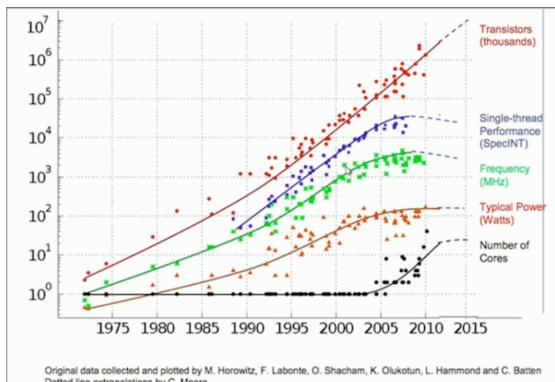
13

- Responsiveness/availability
  - Threads allow a program to continue running even if a part of it is blocked for I/O or is performing a long operation.
- Resource sharing
  - Threads share the memory and other resources of a process that they belong to.
- Efficiency
  - Threads are more lightweight than processes.
    - Process creation is expensive.
    - Switching processes is heavyweight.

14

## Other Viewpoints to Threads

- CPU speed does not increase any more as it did in the past...
  - Physical material barrier to increase clock speed and # of transistors in a CPU
  - Heat (and cooling) and power consumption problems



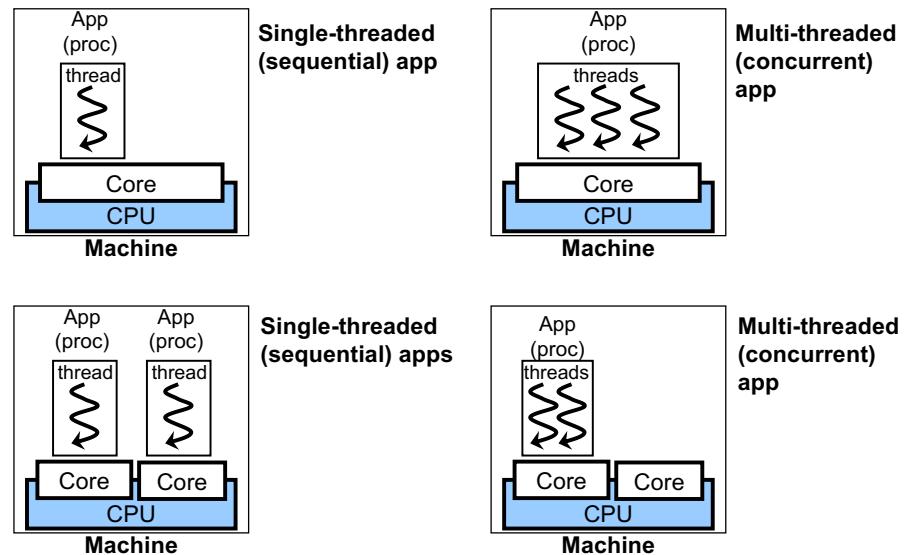
15

- Each CPU increases its density of cores, not its speed.
  - Intel Core i7: up to 4.0 GHz and up to 8 cores
  - Intel Xeon: up to 4.0 GHz and up to 8 cores
  - Multiple relatively-slower CPU cores for mobile devices to reduce power consumption
    - Amazon Kindle Fire (\$50): Quad-core 1.3 GHz
    - Apple iPhone 8/10: 6 cores
- Significant energy consumption and CO<sub>2</sub> emission by data centers
  - Amazon (450K+ servers), Facebook, Google (500K+ servers), Microsoft, Salesforce.com and Yahoo!
  - DCs consumed approx. 1.5% of global electricity usage ('10).
  - Google paid \$2M/mo for electricity bills ('07).
  - The Info. and Comm. Technology (ICT) industry produced 2% of global CO<sub>2</sub> emission ('07).
    - On par with the aviation industry.
  - IDCs were responsible for 23% of the ICT industry's emission

16

## Terminology: Single- or Multi-threaded?

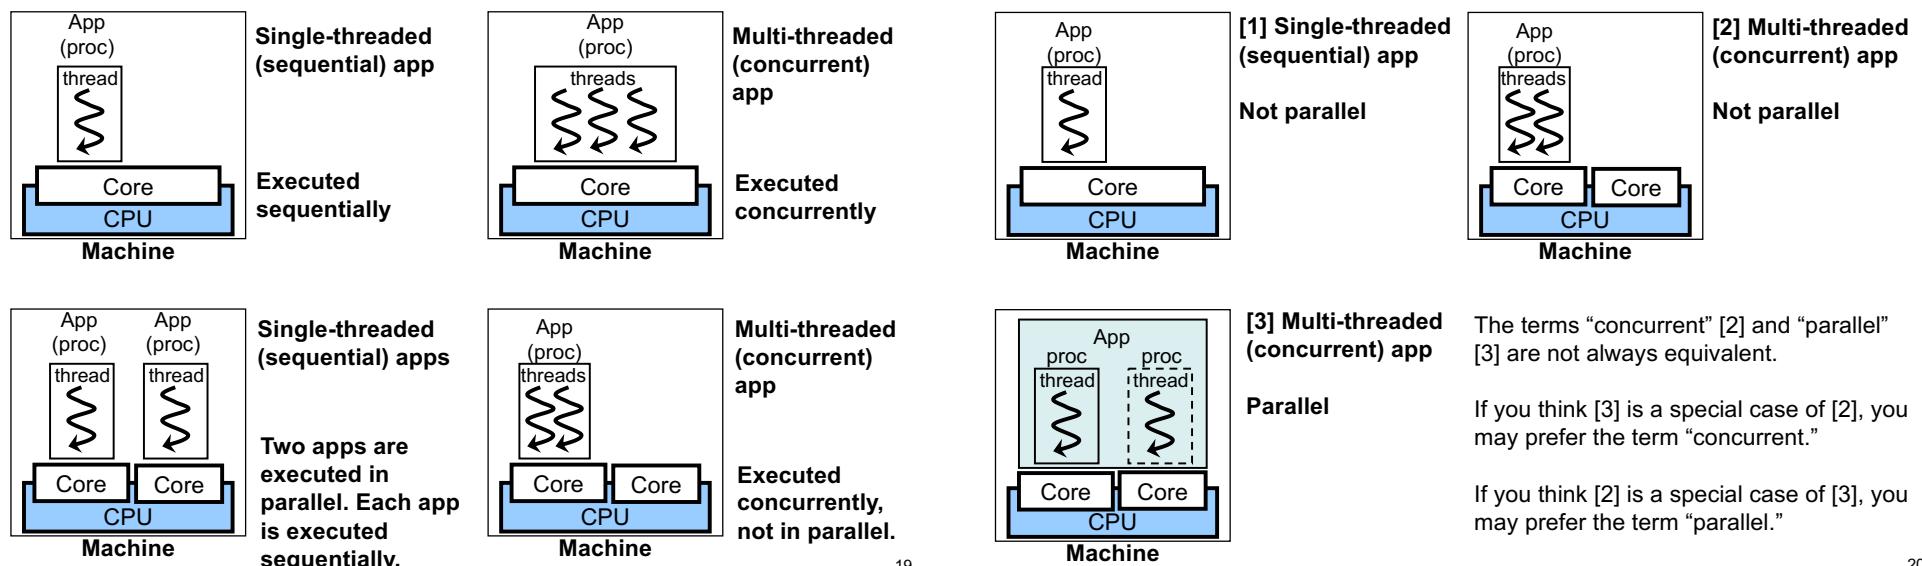
- GPUs are designed with a massive # of cores.
  - NVIDIA Tesla K80: 4992 cores
- You are expected to increase your app's performance by increasing its concurrency.
- Ultimate goals of multi-threaded applications:  
**Responsiveness and performance improvement**



17

18

## Terminology: Concurrent or Parallel?



19

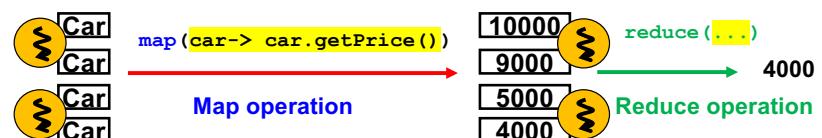
20

## In Java...

- The terms “concurrent” and “parallel” are somewhat mixed up in Java.
  - Java 5 (2004) introduced `java.util.concurrent`, which was further enhanced by Java 6 (2006).
    - “Concurrent” collections such as `ConcurrentHashMap`
    - The Executor framework
      - An extension/abstraction over low-level threads
  - Java 7 (2011)
    - The Fork/Join framework, an extension to the Executor framework
      - The term “parallel” appeared in its documentation, although it was placed in `java.util.concurrent`.
    - “Parallel” garbage collector, which performs garbage collection with multiple threads.
  - Java 8 (2014)
    - “Parallel” streams, which allows for “parallel” (threaded) operations on collections with lambda expressions

- An example of using a **parallel stream**.

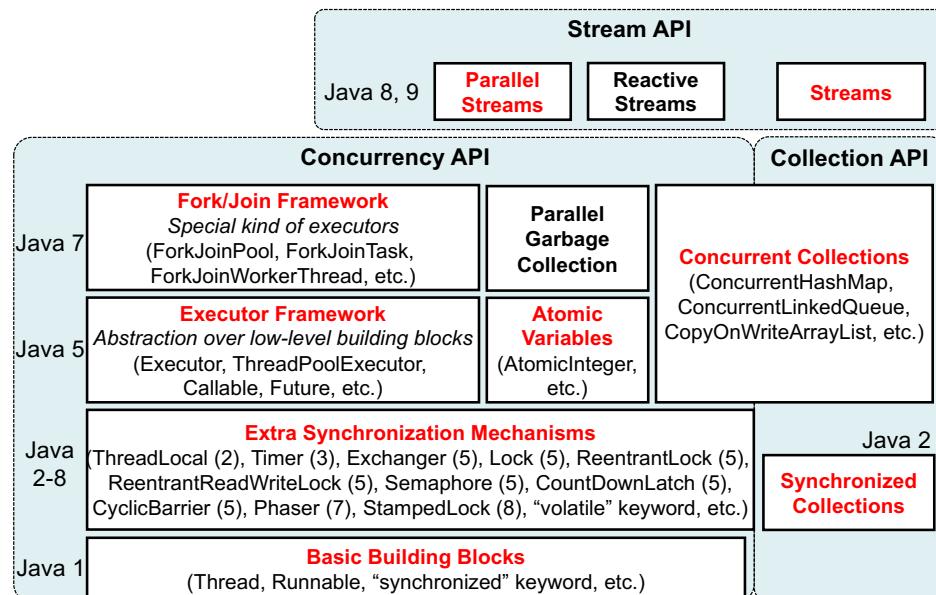
```
- Integer price = cars.stream()
    .parallel()
    .map((Car car) -> car.getPrice())
    .reduce(0, (result, carPrice) ->{
        if(result==0) return carPrice;
        else if(carPrice < result) return carPrice;
        else return result; } );
```



21

22

## Concurrency API in Java



## Java Threads

24

# Java Threads

- Every Java program has at least one *thread of control*.
  - `main()` runs with a thread on a JVM.

- When a JVM starts, it implicitly (or automatically) creates the “main” thread.

```
• class HelloWorldApp{
    public static void main(String[] args){
        System.out.println("Hello World!");
    }
}

• >java HelloWorldApp
```

- If you need **extra threads** in addition to the main thread, you need to *explicitly* create them.
- 4 things to do:
  - Define a class implementing the `java.lang.Runnable` interface
    - `public abstract void run();`
  - Write a threaded/concurrent task in `run()` in the class
  - Instantiate `java.lang.Thread` and associate a Runnable object with the thread
  - Start (call `start()` on) the instantiated thread.
    - `run()` is automatically called on the thread.

25

26

## An Example Code: Creating a Thread

- GreetingRunnable.java

```
- class GreetingRunnable implements Runnable{
    private String greeting;

    public GreetingRunnable(String aGreeting){
        greeting = aGreeting;
    }

    public void run(){
        for( int i=0; i<10; i++ ){
            Date now = new Date();
            System.out.println(now + " " + greeting);
        } } }
```

- HelloWorldTest.java

```
- main(...){
    GreetingRunnable runnable = new GreetingRunnable("Hello World");
    Thread t = new Thread(runnable);
    t.start(); }
```

## Thread.start()

- Creating a Thread object does not mean creating a new thread.

- It is `start()` that actually creates a thread.

- `start()`

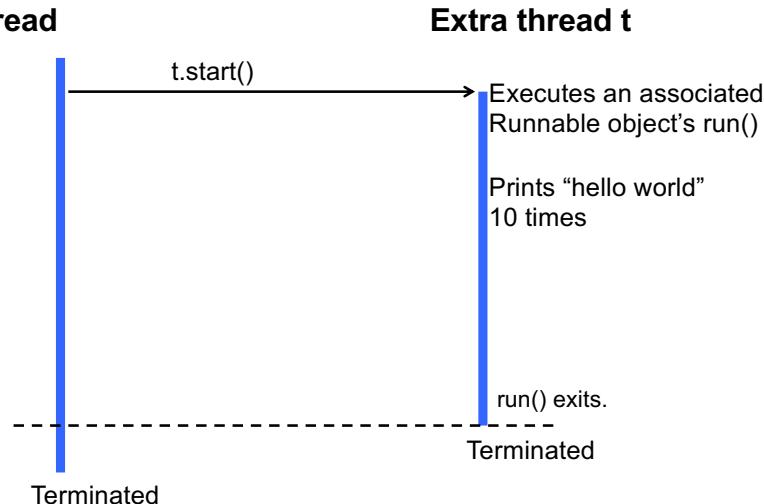
- Allocates memory and initializes a new thread on a JVM.
  - Calls `run()` of a specified Runnable object.
  - Do not call `run()` directly, but let `start()` call `run()` on behalf of yourself.

27

28

# Program Execution

Main thread



- Output:

- Mon Mar 26 15:14:43 EDT 2007 Hello World
- Mon Mar 26 15:14:44 EDT 2007 Hello World
- Mon Mar 26 15:14:45 EDT 2007 Hello World
- Mon Mar 26 15:14:46 EDT 2007 Hello World
- Mon Mar 26 15:14:47 EDT 2007 Hello World
- Mon Mar 26 15:14:48 EDT 2007 Hello World
- Mon Mar 26 15:14:49 EDT 2007 Hello World
- Mon Mar 26 15:14:50 EDT 2007 Hello World
- Mon Mar 26 15:14:51 EDT 2007 Hello World
- Mon Mar 26 15:14:52 EDT 2007 Hello World

29

30

## An Example Code: Creating Threads

- HelloWorldTest2.java and GreetingRunnable.java

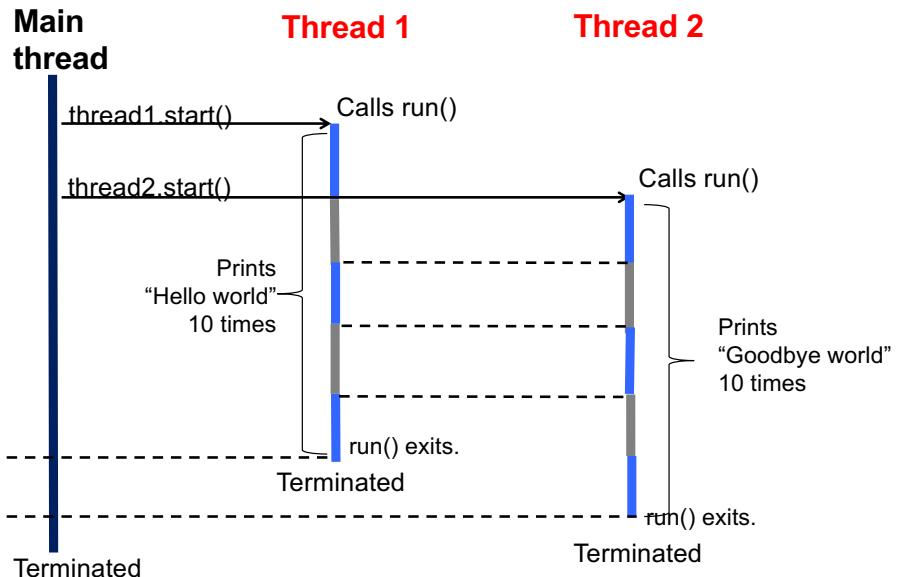
```
main(...){
    GreetingRunnable runnable1 = new GreetingRunnable("Hello World");
    GreetingRunnable runnable2 = new GreetingRunnable("Goodbye World");

    Thread thread1 = new Thread(runnable1);
    Thread thread2 = new Thread(runnable2);

    thread1.start();
    thread2.start();
}
```

## Expected Program Execution

Main  
thread



31

32

## A Possible Output

- Output:
  - Mon Mar 26 15:28:45 EDT 2007 Goodbye World
  - Mon Mar 26 15:28:45 EDT 2007 Hello World
  - Mon Mar 26 15:28:46 EDT 2007 Hello World
  - Mon Mar 26 15:28:46 EDT 2007 Goodbye World
  - Mon Mar 26 15:28:47 EDT 2007 Goodbye World
  - Mon Mar 26 15:28:48 EDT 2007 Goodbye World
  - Mon Mar 26 15:28:48 EDT 2007 Hello World
  - Mon Mar 26 15:28:49 EDT 2007 Goodbye World
  - Mon Mar 26 15:28:49 EDT 2007 Hello World
  - Mon Mar 26 15:28:50 EDT 2007 Goodbye World
  - Mon Mar 26 15:28:50 EDT 2007 Hello World
  - Mon Mar 26 15:28:51 EDT 2007 Goodbye World
  - Mon Mar 26 15:28:51 EDT 2007 Hello World
  - Mon Mar 26 15:28:52 EDT 2007 Hello World
  - Mon Mar 26 15:28:52 EDT 2007 Goodbye World
  - Mon Mar 26 15:28:53 EDT 2007 Hello World
- Two message sets (Hello and Goodbye) are **NOT exactly interleaved.**

33

## The Order of Thread Execution

- JVM's thread scheduler gives you NO guarantee about the order of thread execution.
- There are always slight variations in the time to execute a threaded task
  - especially when calling OS system calls (typically I/O related system calls)
- Expect that the order of thread execution is somewhat random.

34

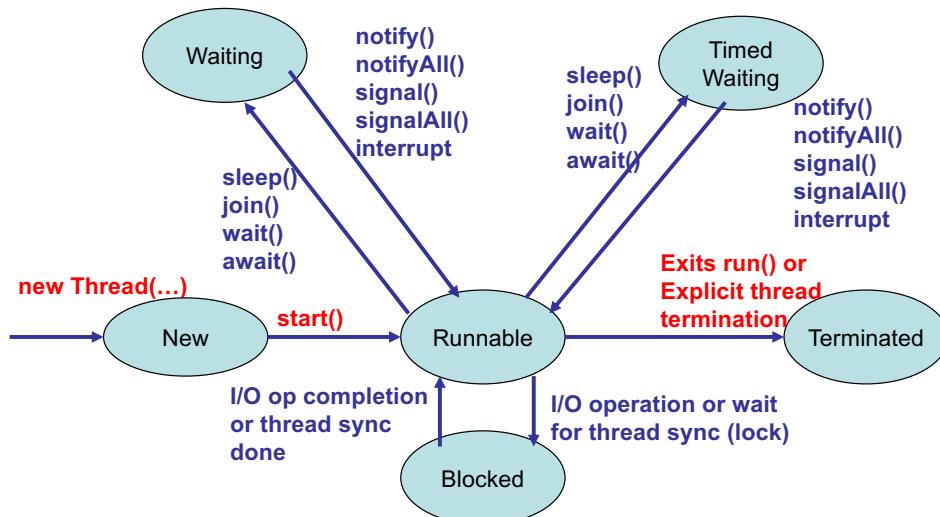
## Exercise

- Modify HelloWorldTest2.java;
  - Replace the following 4 lines
    - Thread thread1 = **new** Thread(runnable1);
    - Thread thread2 = **new** Thread(runnable2);
    - thread1.start();
    - thread2.start();
  - with the following 2 lines:
    - runnable1.run();
    - runnable2.run();
- What does the program output?
- Modify HelloWorldTest2.java;
  - Replace the following 4 lines
    - Thread thread1 = **new** Thread(runnable1);
    - Thread thread2 = **new** Thread(runnable2);
    - thread1.start();
    - thread2.start();
  - with the following 2 lines:
    - runnable1.run();
    - runnable2.run();
- Output: “Hello World” in 10 lines, followed by “Good Bye World” in 10 lines.
  - The program runs **sequentially**, not concurrently.

35

36

# States of a Thread



37

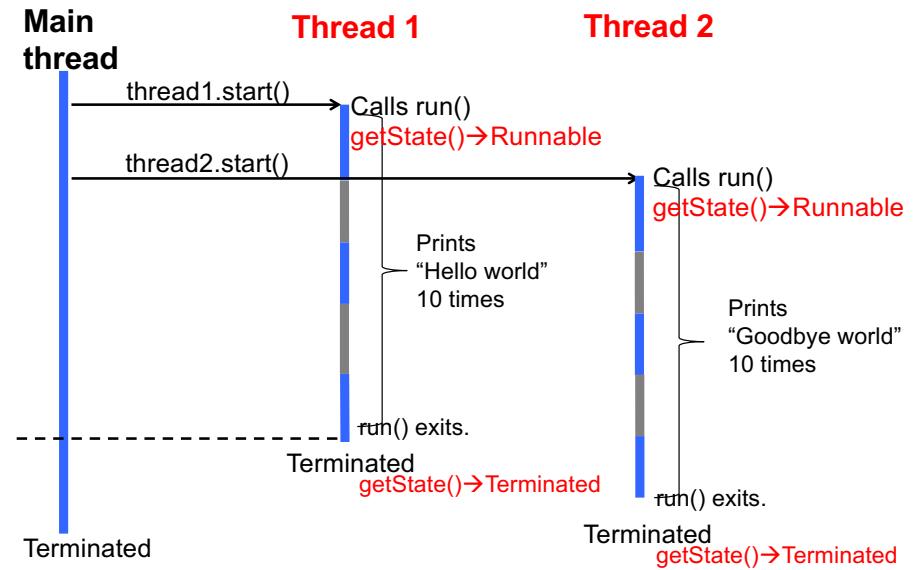
- New
  - A Thread object is created. `start()` has not been called on the object yet.
- Runnable
  - Java does not distinguish **ready-to-run** and **running**. A running thread is still in the Runnable state.
- Terminated/dead
  - A thread automatically dies when `run()` returns.

38

## Program Execution w/ HelloWorldTest2

- ```
public class Thread{
    public enum State{
        NEW, RUNNABLE, BLOCKED, WAITING,
        TIMED_WAITING, TERMINATED
    }

    public Thread.State getState()
    public boolean isAlive()
}
```
- Alive
  - in the Runnable, Blocked, Waiting or Timed Waiting state.
  - `isAlive()` is usually used to poll/check a thread to see if it is terminated.



39

40