




Exercise: Concurrent Access to a File

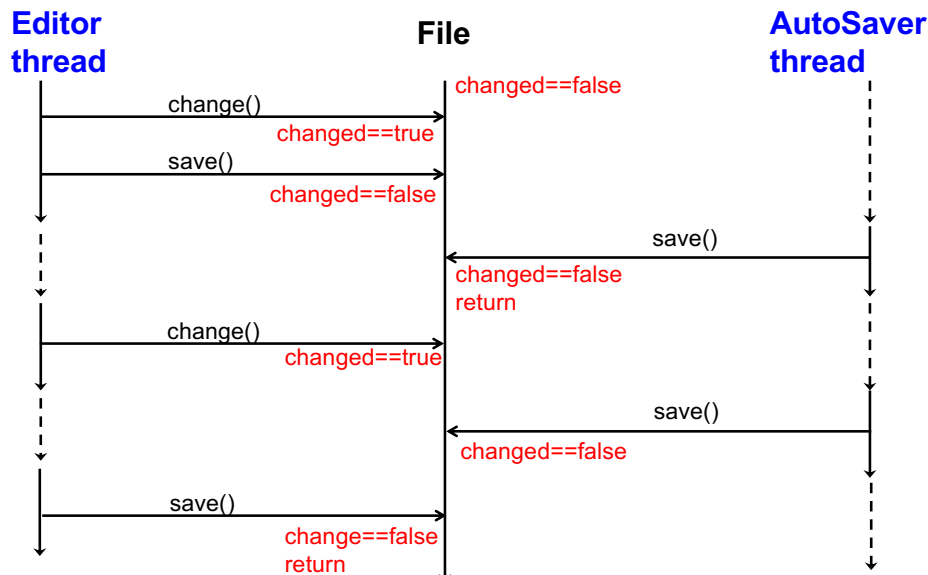


- Imagine a word processing app.
- Assume two threads
 - One for **editing a file**
 - Allows the user to edit a file and save it.
 - One for **saving a file automatically**
 - Periodically saves an open file at background.
- The 2 threads call `change()` and `save()` on an open file concurrently.



- File**
 - Has a boolean variable: "changed"
 - Initialized to be false.
 - `change()`
 - Changes the file's content.
 - Assigns true to the variable "changed."
 - `save()`
 - if(!changed) return;
 - if(changed)
 - Print out some message (e.g., time stamp, etc.)
 - Assign false to the variable "changed."
- Editor (a Runnable) repeats:**
 - Calls `change()` and `save()`
 - Sleeps for a second.
- AutoSaver (a Runnable) repeats:**
 - Calls `save()`
 - Sleeps for two seconds.

Desirable Result



HW 7

- Race conditions can occur if you do not guard the variable `changed` with a lock. Explain a potential race condition with a diagram like in a previous slide.
- Implement `File`, `Editor` and `AutoSaver` in a thread-safe manner
 - Define a `ReentrantLock` in `File`. Use the lock in `change()` and `save()`
 - c.f. `deposit()` and `withdraw()`, which use a lock to access a shared variable in the bank account example
 - Use try-finally blocks: Always do this in all subsequent HWs.
 - Create two extra threads and have them execute `Editor's run()` and `AutoSaver's run()`
 - Those threads acquire and release the lock in `change()` and `save()`

- Implement explicit thread termination in `Editor` and `AutoSaver` to terminate 2 extra threads.
- Have the main thread sleep for some time while `Editor` and `AutoSaver` are running.
 - Use `Thread.sleep()`
- Have the main thread terminate the two threads.
 - Define a flag variable `done` and `setDone()` in `Editor` and `AutoSaver`
- Note that this sample code is not thread-safe.
 - Define a `ReentrantLock` in each of `Editor` and `AutoSaver` to guard a flag variable `done`.
 - Use try-finally blocks
 - Use balking in `run()`
 - Do not surround a "while" loop with `lock()` and `unlock()`.

```
class Editor implements Runnable{
    private boolean done = false;

    public void run(){
        while(true){
            if(done){
                System.out.println("...");
                break;
            }
            aFile.change();
            aFile.save();
            Thread.sleep(1000);
        }
    }

    public void setDone(){
        done = true; } }



```

- Deadline: April 2 (Tue) midnight

Recap: Singleton Design Pattern


- Guarantee that a class has only one instance.
 - c.f. CS680 lecture note
- ```
public class Singleton{
 private Singleton(){};
 private static Singleton instance = null;

 // Factory method to return the singleton instance
 public static Singleton getInstance(){
 if(instance==null)
 instance = new Singleton();
 return instance;
 }
}
```
- This code is NOT thread-safe; race conditions can occur.



- When multiple threads call `getInstance()` concurrently, they share instance.
- ```
public class Singleton{
    private Singleton(){};
    private static Singleton instance = null;

    // Factory method to return the singleton instance
    public static Singleton getInstance(){
        if(instance==null)           // Read   3 steps
            instance = new Singleton(); // Write  2 steps
        return instance;             // Read   2 steps
    }
}
```



- ```
public class Singleton{
 private Singleton(){};
 private static Singleton instance = null;
 // Write logic. Requires 2 steps
 // Thread safe }
}
```
- JVM completes all **initial value assignments on all static data fields** BEFORE using a class or creating class instances.
  - `instance` has been initialized before a thread(s) call `getInstance()`
  - This write logic is thread-safe.
    - No need to worry about race conditions here.

9

## Concurrent Singleton Design Pattern

- Guarantee that a class has only one instance.
- ```
public class ConcurrentSingleton{
    private Singleton(){};
    private static Singleton instance = null;
    private static ReentrantLock lock = new ReentrantLock();

    // Factory method to create or return the singleton instance
    public static Singleton getInstance(){
        lock.lock();
        try{
            if(instance==null){ instance = new Singleton(); }
            return instance;
        }finally{
            lock.unlock();
        }
    }
}
```

10

HW 8

- The Singleton class is not thread-safe.
 - Race conditions can occur if you do not guard the `instance` variable with a lock. Explain a potential race condition in which more than one instances are created.
 - Use a diagram like in a previous slide.
- Submit a thread-safe version of it (`ConcurrentSingleton`)
 - Define a lock in Singleton. Use the lock in `getInstance()`
 - Use try-finally blocks: Always do this in all subsequent HWs.
 - Create multiple extra threads and have them call `getInstance()`
 - Make sure that only one instance is created.
 - Use `System.out.println(Singleton.getInstance())`

- Deadline: April 2 (Tue) midnight

Regular and Static Locks

- ```
public class Foo{
 ReentrantLock lock = new ReentrantLock();
 static ReentrantLock sLock = new ReentrantLock(); }
```
- A **regular lock** is created and used on an instance-by-instance basis.
  - Different instances of `Foo` have **different** locks (i.e., different instances of `ReentrantLock`).
- A **static lock** is created and used on a per-class basis.
  - All instances of `Foo` share a **single** lock (`sLock`).

- ```
public class Foo{
    private ReentrantLock lock = new ReentrantLock();
    private static ReentrantLock sLock = new ReentrantLock();

    public void a() {...}
    public void b() {...}
    public void syncA() {lock.lock(); ... lock.unlock();}
    public void syncB() {lock.lock(); ... lock.unlock();}

    public static void sA() {...}
    public static void sB() {...}
    public static void sSyncA() {sLock.lock(); ... sLock.unlock();}
    public static void sSyncB() {sLock.lock(); ... sLock.unlock();} }
```
- `x = new Foo(); y = new Foo();`
- Two threads call...
 - `x.a()` and `Foo.sA()`: No synchronization for the two threads
 - `x.syncA()` and `Foo.sA()`: No synchronization
 - `Foo.sA()` and `Foo.sA()`: No synchronization
 - `Foo.sA()` and `Foo.sB()`: No synchronization
 - `x.syncA()` and `Foo.sSyncA()`: No synchronization
 - `Foo.sSyncA()` and `Foo.sSyncA()`: Synchronization
 - `Foo.sSyncA()` and `Foo.sSyncB()`: Synchronization
 - `x.sSyncA()` and `y.sSyncB()`: Synchronization
 - This is not grammatically wrong, but **write `Foo.sSyncA()` instead of `x.sSyncA()`**

Exercise: Regular and Static Locks

- ```
public class Foo{
 private ReentrantLock lock = new ReentrantLock();

 public void a() {...}
 public void b() {...}
 public void syncA() {lock.lock(); ... lock.unlock();}
 public void syncB() {lock.lock(); ... lock.unlock();}
}
```
- `x = new Foo(); y = new Foo();`
- Two threads call...
  - `x.a()` and `x.a()`: no synchronization (no mutual exclusion) for the two threads
  - `x.a()` and `x.b()`: no synchronization
  - `x.a()` and `x.syncA()`: no synchronization
  - `x.syncA()` and `x.syncA()`: Synchronization (mutual exclusion)
  - `y.syncA()` and `y.syncB()`: Synchronization
  - `x.syncA()` and `y.syncA()`: No synchronization
  - `x.syncA()` and `y.syncB()`: No synchronization

## Thread.sleep()


- ```
Thread t = new Thread( new FooRunnable() );
t.start();
try{
    t.sleep(1000);
}catch(InterruptedException e){...}
```
- It looks like an extra thread (`t`) will sleep.
- However, the main thread will actually sleep
 - because `sleep()` is a **static method** of `Thread`.
 - `Thread.sleep()`: Allows the **currently executing thread** to sleep (temporarily cease execution) for the specified number of milliseconds
- DO NOT write `t.sleep(...)`. It's misleading and error-prone.
- ALWAYS WRITE `Thread.sleep(...)`.
 - Make sure to do this in HW 7.

RunnableInterruptiblePrimeGenerator

```
class InterruptiblePrimeGenerator extends PrimeGenerator {
    public void generatePrimes() {
        for (long n = from; n <= to; n++) {
            if ( Thread.interrupted() ) {
                System.out.println("Stopped");
                this.primes.clear();
                break;
            }
            if ( isPrime(n) ) { this.primes.add(n); } } }
}
```

```
class RunnableInterruptiblePrimeGenerator
    extends InterruptiblePrimeGenerator
    implements Runnable {

    public void run() {
        generatePrimes(); } }
}
```



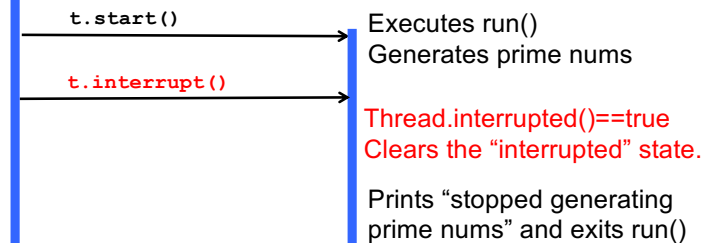
interrupt(), isInterrupted() and interrupted()

- public class Thread{
 public void interrupt();
 public boolean isInterrupted();
 public static boolean interrupted();
 }
- Each thread (Thread instance) has the “interrupted” (boolean) state.
- interrupt()
 - Interrupts **this** thread and changes its “interrupted” state.
 - aThread = new Thread(...); aThread.start();
 aThread.interrupt();
- isInterrupted()
 - Returns true if **this** thread has been interrupted.
 - aThread = new Thread(...); aThread.start();
 if (aThread.isInterrupted()) { ... }
 - Does not change the “interrupted” state of the thread.
- interrupted()
 - Returns true if the **currently-executed thread** has been interrupted.
 - Clears the “interrupted” state (true → false) if true is returned.

Main thread

Thread t

```
gen = new RunnableInterruptiblePrimeGenerator(...)
t = new Thread(gen)
```



```
for( long n = from; n <= to; n++ ){
    // Detect if another thread has interrupted. Balk if yes.
    if( Thread.interrupted() ){
        System.out.println("Stopped generating prime nums.");
        this.primes.clear();
        break; // balking
    }
    if( isPrime(n) ) { this.primes.add(n); } } }
```

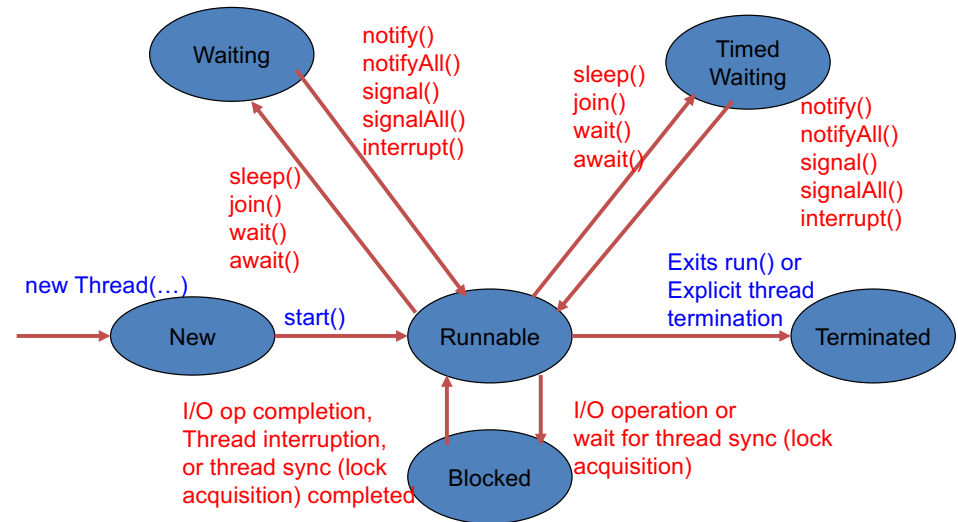
Thread Interruption != Thread Termination

- interrupt() **NEVER** terminate a thread.
 - It simply change the “interrupted” state
 - to help/trigger a thread termination.

What Happens When `interrupt()` is Called on a Thread?

- If the soon-to-be-terminated thread is in the Runnable state, `interrupt()` changes its “interrupted” state to be true.
- If the soon-to-be-terminated thread is in the *Waiting* or *Blocked* state, it throws an `InterruptedException`.

States of a Thread



22

RunnableInterruptiblePrimeGenerator

- In fact, `RunnableInterruptiblePrimeGenerator` is NOT thread-safe. Race conditions can occur.

```

class InterruptiblePrimeGenerator extends PrimeGenerator {
    public void generatePrimes(){
        for (long n = from; n <= to; n++){
            if( Thread.interrupted() ){ // 2 steps
                System.out.println("Stopped");
                this.primes.clear();
                break;
            }
            if( isPrime(n) ){ this.primes.add(n); }
        }
    }
}

class RunnableInterruptiblePrimeGenerator
    extends InterruptiblePrimeGenerator
    implements Runnable {

    public void run(){
        generatePrimes();
    }
}
  
```

Thread.interrupt()

- ```
public void interrupt(){
 ...
 synchronized(...){ // Acquire a lock in Thread
 ...
 interrupt0(); // native method (atomic)
 ...
 }
}
```
- `interrupt()` and `interrupted()` are thread-safe.
  - `isInterrupted()` is thread-safe as well.
  - c.f. Java source code
- However, *client code* of `interrupted()` is NOT guaranteed to be thread-safe.

# Solution: Locking and Balking

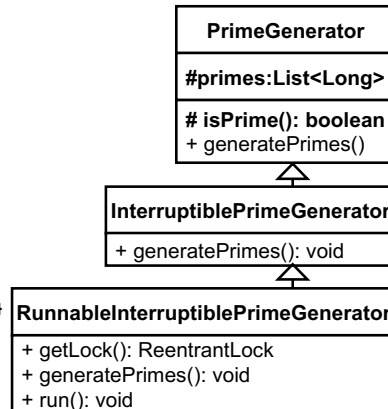
```
class RunnableInterruptiblePrimeGenerator
 extends InterruptiblePrimeGenerator
 implements Runnable {

 private final ReentrantLock lock = new ReentrantLock();

 public ReentrantLock getLock() {
 return lock; }

 public void generatePrimes() {
 for (long n = from; n <= to; n++) {
 lock.lock();
 if(Thread.interrupted()){
 System.out.println("Stopped");
 this.primes.clear();
 break;
 }
 lock.unlock();
 if(isPrime(n)) {this.primes.add(n);}}

 public void run(){
 generatePrimes(); } }
```



## HW 9

- Revise `RunnableInterruptiblePrimeGenerator.java` to be thread-safe.
  - c.f. HW 6, in which you work on a thread-safe version of `RunnableCancelablePrimeGenerator.java`
- Deadline: April 2 (Tue) midnight

- Main thread (client of `RunnableInterruptiblePrimeGenerator`)

```
- RunnableInterruptiblePrimeGenerator gen =
 new RunnableInterruptiblePrimeGenerator();
 Thread aThread = new Thread(gen); aThread.start();
```

```
gen.getLock().lock();
aThread.interrupt();
gen.getLock().unlock();
```

- This code uses two locks.
  - One in `Thread`
  - One in `RunnableInterruptiblePrimeGenerator`

## Explicit Thread Termination

- Flag-based
  - Pros:
    - Uses **1 lock (faster)**
  - Cons:
    - Program responsiveness may be lower.
      - if a flag-flipping (e.g. `done==false → true`) happens when a soon-to-be-terminated thread is in the Waiting or Blocked state.
- Interruption-based
  - Pros
    - Higher program responsiveness
      - `interrupt()` can immediately wake up a soon-to-be-terminated thread that is in the Waiting or Blocked state
  - Cons
    - Uses **2 locks (slower)**

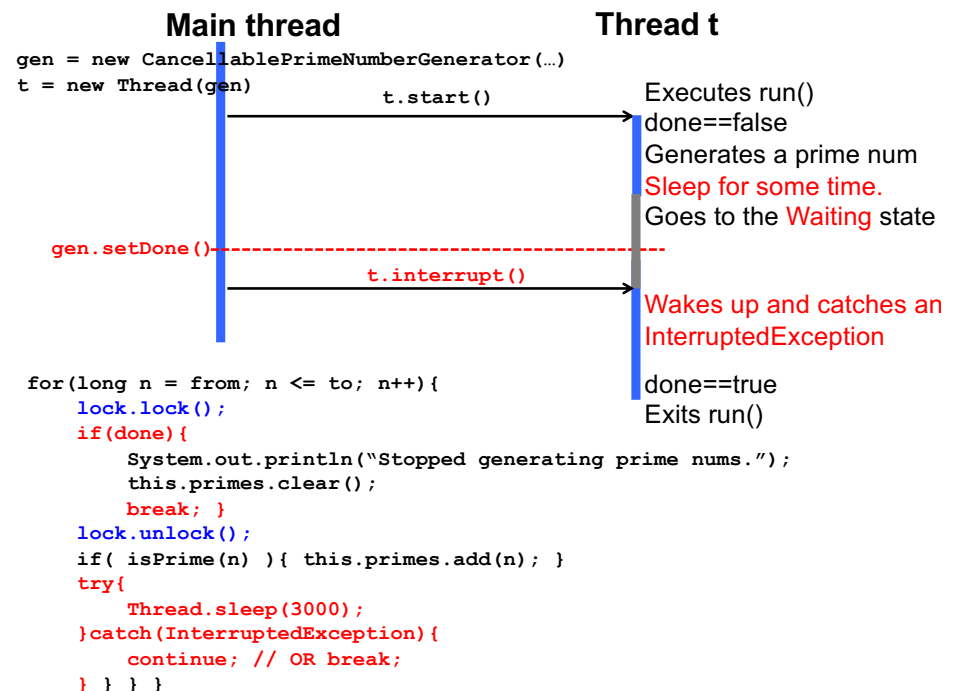
## Hybridization of the Two Approaches?

- Can we implement a **responsive** thread termination that uses **only 1 lock**?

## 2-Step Thread Termination ("Graceful" Thread Termination)

### 2-Step Thread Termination

- Primarily takes the flag-based approach.
  - A soon-to-be-terminated thread periodically checks a flag.
- Let the "terminator" thread call `interrupt()` after flipping the flag's state (i.e., after calling `setDone()`)



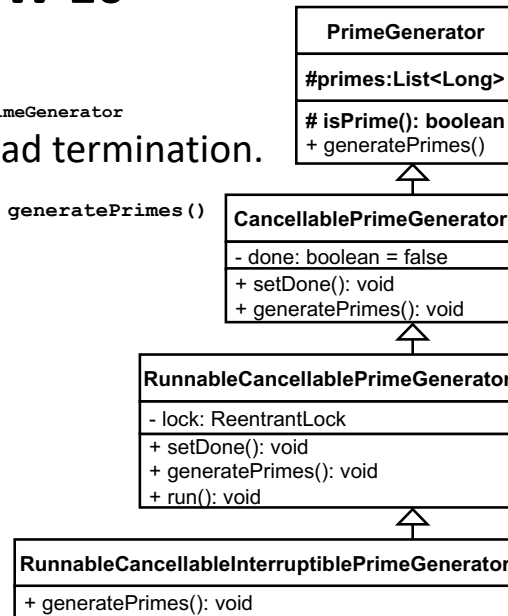


## HW 10

- Define `RunnableCancellableInterruptiblePrimeGenerator` to perform 2-step thread termination.

– Re-define (or override) `generatePrimes()`

- Deadline: April 4 (Thu) midnight



## 2-Step Thread Termination is Effective if...

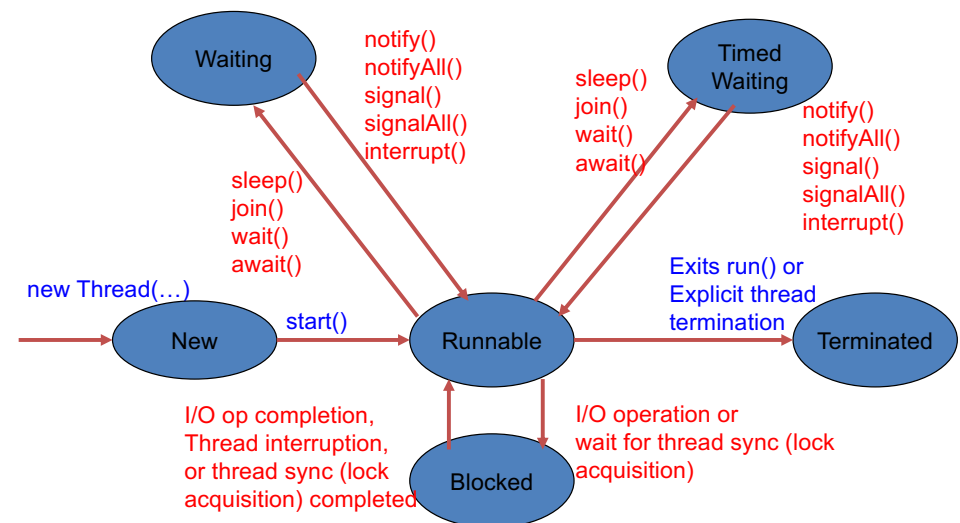
- A soon-to-be-terminated thread may be in the **Waiting or Blocked state** when a “terminator” thread tries to terminate it.
  - Performing an I/O operation.
    - e.g., reading/writing data from/to a file, waiting for an incoming data on a socket, sending data to a remote app.
  - Waiting for a lock acquisition
    - Has called `lock()` on a lock, but the lock is not available yet.
  - Has called `sleep()`, `join()`, etc.

## What Happens

### When `interrupt()` is Called on a Thread?

- If a soon-to-be-terminated thread is in the `Runnable` state, `interrupt()` changes its “interrupted” state to be true.
- If the soon-to-be-terminated thread is in the *Waiting or Blocked* state, it raises an `InterruptedException`.

## States of a Thread



# InterruptedException

- Some methods in Java API throws

`InterruptedException`.


- They can respond to a thread interruption by throwing an `InterruptedException`.

- `Thread.sleep()`
- `Thread.join()`
- `ReentrantLock.lockInterruptibly()`
- `BlockingQueue.put()/take()`
- `Condition.await()`
- I/O operations

- These methods can be long-running and **interruptible**.

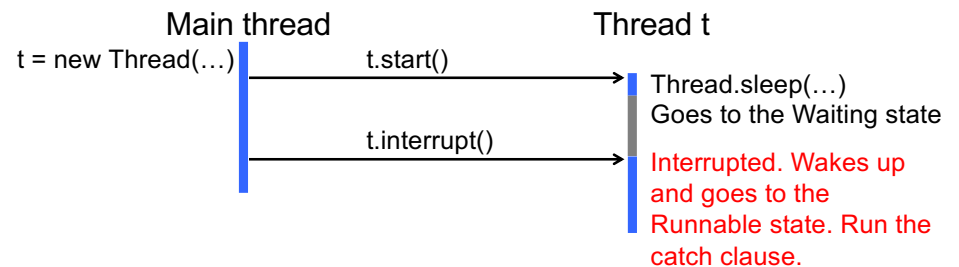
## Thread.join()

- `join()` lets the *currently-executed thread* to wait/sleep until another thread terminates (i.e., until another thread returns `run()`).
- `interrupt()` can interrupt a waiting/sleeping thread.
  - Force `join()` to throw an `InterruptedException`.




## Thread.sleep()

- `sleep()` lets the *currently-executed thread* to sleep for a specified time period.
- `interrupt()` interrupts a sleeping thread.
  - Wakes up the thread and force `sleep()` to throw an `InterruptedException`.
- ```
try{
    Thread.sleep(60000);
}catch(InterruptedException e){
    // Write thread termination (shutdown) logic here.
}
```



Thread.join()

- `join()` lets the *currently-executed thread* to wait/sleep until another thread terminates (i.e., until another thread returns `run()`).
- `interrupt()` can interrupt a waiting/sleeping thread.
 - Force `join()` to throw an `InterruptedException`.




Condition.await()

- `await()` lets the currently-executed thread wait/sleep until another thread wakes it up with `signal()/signalAll()`.
- `interrupt()` can interrupt a waiting/sleeping thread.
 - Allows `await()` to acquire a lock and forces it to throw an `InterruptedException`

```
withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        try{
            // waiting for the balance to exceed 0
            sufficientFundsCondition.await();
        }catch(InterruptedException e){
            //Do something
        }
    }
    belowUpperLimitFundsCondition.signalAll();
    balance -= amount;
    lock.unlock();
}
```

Main thread

Withdraw thread




```
withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        try{
            // waiting for the balance to exceed 0
            sufficientFundsCondition.await();
        }catch(InterruptedException e){
            //Do something; e.g., balk with a
            // "break" statement.
        }
    }
    ...
}
```

BlockingQueue

- `interface BlockingQueue<E> extends Queue<E>`
 - Adds A Queue that additionally supports operations that
 - wait for the queue to become non-empty when retrieving an element
 - wait for space to become available in the queue when storing an element.
- Several impls: `ArrayBlockingQueue`, `LinkedBlockingQueue`, etc.
 - `put()` and `take()` are *blocking* methods.
 - `put()`: Add an element to a queue as the last element.
 - `take()`: Get the first element in the queue.
 - They can respond to a thread interruption by throwing an `InterruptedException`.

Thread t



Thread Termination

- Thread creation is a no brainer.
- Thread termination requires your careful attention.
 - No methods available in `Thread` to directly terminate threads like `terminate()`.
 - Do: 2-step termination
 - Why not?
 - Different programmers/apps need different termination policies.
 - Notify on-going thread termination to other threads?
 - Raise exception(s) in addition to `InterruptedException`?
 - What to do for the data maintained by a thread being terminated?
 - Java allows you to flexibly craft your own termination policy.

Where did the `synchronized` Keyword go?

- Java still has the `synchronized` keyword.

- ```
public synchronized void foo(){
 // The entire method body is atomic.
}
```
- ```
public void foo(){  
    // non-atomic code here  
    synchronized(this){  
        // atomic code here  
    }  
    // non-atomic code here }
```

- Implicit locking

- No need to create a `ReentrantLock` and call `lock()` and `unlock()`.
- When a thread enters a synchronized method/block, it tries to acquire the (implicit) lock that `this` instance maintains.
 - Instance-by-instance locking
- Code gets tricky/dirty to use multiple locks in a single class.

45

- Explicit locking

- ```
ReentrantLock aLock = new ReentrantLock()
public void foo(){
 aLock.lock();
 // atomic code
 aLock.unlock(); }
```

- Arbitrary locking scope.

- Clean code even if a class uses multiple locks.

- Extra functionalities

- e.g., `getQueueLength()`: returns the # of waiting threads.
- `tryLock()`: acquires a lock only if it is not held by another thread.

- The catch is... it's **VERY easy to forget calling `unlock()`**.

- Must call `unlock()` in a finally clause.

- Implicit locking with the “synchronized” keyword

- A thread can call `notify()` and `notifyAll()` even if it has not acquired a lock.
  - An `IllegalMonitorStateException` is thrown.

- Explicit locking

- This error/bug never occurs.

- ```
ReentrantLock lock = new ReentrantLock();  
Condition cond = lock.newCondition();  
lock.lock();  
...  
cond.signalAll();
```