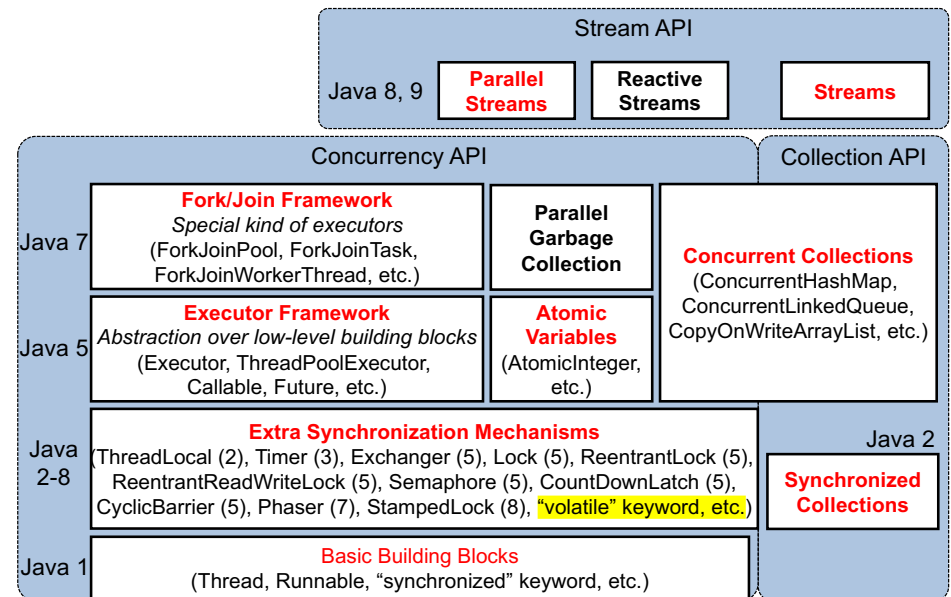


Volatile Variables in Java

Concurrency API in Java



What is the “volatile” Keyword?

- You must do thread synchronization (locking) to share a variable among threads.
- You can skip locking to share a **volatile variable** among threads in some cases.
 - No locking means...
 - No worry about thread safety issues
 - Particularly, race conditions
 - Less overhead (better performance)
 - Locking consumes some time and resources.

Recap: Race Condition

- Explicit thread termination
- A thread (“A”) has read the current value of “done” but has not checked if it is true or false.
- A context switch occurs.
- A race condition occurs if
 - Another thread (“B”) calls setDone() to assign true to “done.”
 - This state update is not visible for Thread “A”.
 - It still believes done==false and does not immediately break out from the loop.

```
boolean done = false;

public void setDone(){
    done = true;
}

public void run(){
    while( true )
        if( done ) break;
        // Do some work
}
```

Thread-safe Code

- Without “volatile”

```
boolean done = false;
ReentrantLock lock =
    new ReentrantLock();

public void setDone(){
    lock.lock();
    done = true;
    lock.unlock();
}

public void run(){
    while( true ){
        lock.lock();
        if( done ) break;
        // Do some work
        lock.unlock();
    }
}
```

- With “volatile”

```
volatile boolean done = false;

public void setDone(){
    done = true;
}

public void run(){
    while( true )
        if( done ) break;
        // Do some work
}
```

- Both versions are thread-safe.
- “volatile” makes code simpler and less error-prone.

5

- With “volatile”

```
volatile boolean done = false;

public void setDone(){
    done = true;
}

public void run(){
    while( true )
        if( done ) break;
        // Do some work
}
```

- A context switch can occur in b/w
 - Reading the value of “done”
 - Judging if it is true or not.
- The most up-to-date value of “done” is loaded from the memory in the 2nd step
 - In case another thread has called setDone() in between the 2 steps
- No race conditions occur.
 - “Volatile” makes a state change on a shared variable visible for threads.

Limited Effectiveness/Usefulness

- With “volatile”

```
volatile boolean done = false;

public void setDone(){
    done = true;
}

public void run(){
    while( true )
        if( done ) break;
        // Do some work
}
```

- A context switch can occur in between
 - Loading the true value
 - Assigning it to “done”
- All threads will assign true to “done.”
 - There are no other possible state changes.
- Threads do not generate race conditions.

- A “volatile” variable is guaranteed to have the most up-to-date value whenever it is read.

```
- volatile int a;
int b = a;
if(a==0){...}
println(a);
// These multi-step ops are all thread-safe,
// even if a context switch occurs in between
// 2 different steps and another thread changes
// the value of “a” there.
```

- However, it does NOT eliminate all possible race conditions.

```
- a + 1;
b = a + 1;
if(a+1>0)
println(a+1);
// 3 steps. Thread-safe.
// These operations are NOT thread-safe.
// The first 3 steps are thread-safe though.
// The 3rd step generates an intermediate
// state. Race conditions can occur in b/w the
// 3rd and the 4th step.
```

- “volatile” is effective only in the read operations that have no intermediate state.

- “volatile” is effective only in the write operations that have **no intermediate state**.

```
- volatile int a;
  a = 1;           // 2 steps. A context switch can occur in
                  // between the 2 steps, but no race conditions
                  // occur here.

  a = a + 1;       // 5 steps. There is an intermediate state.
  a++;            // Not thread-safe.
```

- Use a “volatile” variable only when you can live with these limitations/constraints.
- Do NOT use “volatile” for arrays.

In Summary...

- “Volatile” is **NOT a general-purpose, widely-applicable** threading tool.
- Powerful only in some specific cases
 - In practice, assume it is useful only for simplifying the implementation of a latch.
 - **Very useful** to implement **flag-based thread termination** and **2-step thread termination**.

When to use Volatile Variables?

- Despite these limitations, when can we take advantage of volatile variables?
- **ONLY when a value assignment to a shared variable does not depend on the current value.**
- Any data structures that perform this kind of value assignments?
 - Probably, **latch** only.
 - A data structure that performs a single type of state changes
 - e.g. False → True
 - Often used to terminate threads.
 - c.f. “done” variable in prior examples
 - » The state of “done” always changes in a unidirectional way: false → true
 - » “true → false” never happen.

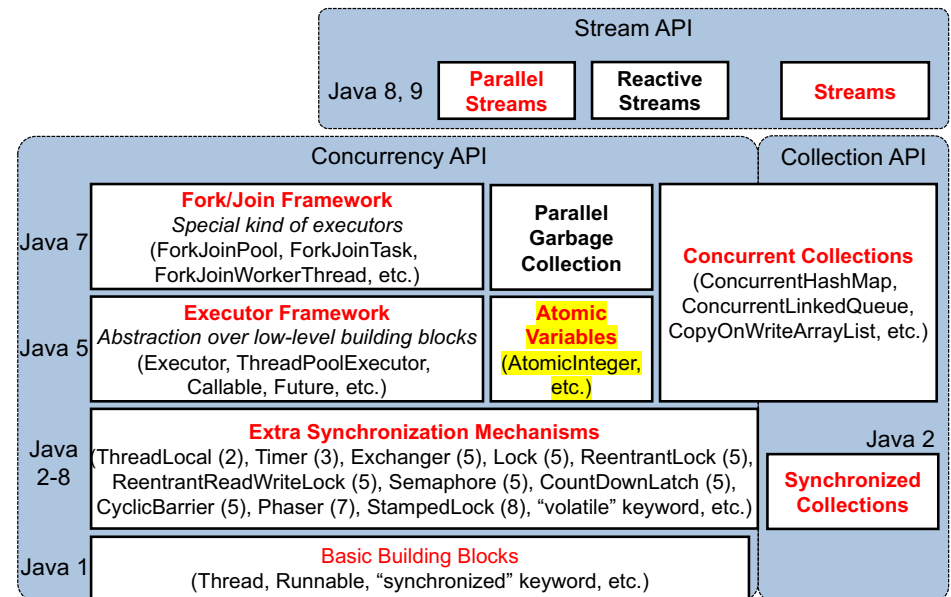
10

Exercise (Not HW)

- Define a flag as a volatile variable in a 2-step thread termination scheme that you have implemented for a prior HW.

Atomic Data Structures in Java

Concurrency API in Java



java.util.concurrent.atomic Package

- Offers **thread-safe classes** to manipulate single variables.
 - AtomicBoolean,
 - AtomicInteger, AtomicIntegerArray
 - AtomicLong, AtomicLongArray
 - AtomicReference<V>, AtomicReferenceArray<E>
 - DoubleAccumulator, DoubleAdder
 - LongAccumulator, LongAdder
 - ...

java.util.concurrent.atomic Package

- Offers **thread-safe classes** to manipulate single variables.
 - AtomicBoolean,
 - AtomicInteger, AtomicIntegerArray
 - AtomicLong, AtomicLongArray
 - AtomicReference<V>, AtomicReferenceArray<E>
 - DoubleAccumulator, DoubleAdder
 - LongAccumulator, LongAdder
 - ...
- All of their public methods are **thread-safe**. They **avoid race conditions**
 - with a special CPU instruction, called Compare-and-Swap (CAS) instruction, rather than doing thread synchronization
 - Very efficient.

Atomic Variables

- Serve as “better” volatile variables.
- Offer the same memory semantics as volatile variables, but with additional support for atomic state updates.
 - `get()` has the memory effects of reading a volatile variable.
 - Returned value is guaranteed to be the most up-to-date whenever it is used.
 - `set()` has the memory effects of writing to a volatile variable.
 - *read-and-update* methods (e.g., `xxxAndSet()` and `getAndXxx()`) have the memory effects of both reading and writing volatile variables.
- Highly recommended as far as they match your use cases.

```
- volatile boolean vFlag;
  vFlag = true;           // 2 steps. Thread-safe.
  if(vFlag){ int i=10; }   // Thread-safe.
  if(vFlag){ vFlag = false; } // NOT thread-safe.
                          // Read and write ops are thread-safe
                          // each, but if a context switch occurs
                          // in between the two ops, a race
                          // condition can occur.

- AtomicBoolean atomicFlag = new AtomicBoolean(false);
  atomicFlag.set(true)      // Thread-safe
  atomicFlag.compareAndSet(true, false); // Thread-safe!
```

- `set()`: Thread-safe. Atomically returns the current value, with memory effects of reading a volatile variable.
- *read-and-update* methods (e.g., `xxxAndSet()` and `getAndXxx()`) have the memory effects of both reading and writing volatile variables.

AtomicBoolean

- Offers thread-safe methods to manipulate a single boolean value atomically (i.e., in a thread-safe manner).

```
- boolean flag;
  if(flag)           // 2 steps. Not thread-safe
  println(flag);     // 2 steps. Not thread-safe

- volatile boolean vFlag;
  if(vFlag)           // 2 steps. Thread-safe
  println(vFlag);     // 2 steps. Thread-safe

- AtomicBoolean atomicFlag = new AtomicBoolean(false);
  if(atomicFlag.get()) // 2 steps. Thread-safe
  println(atomicFlag.get()); // 2 steps. Thread-safe
```

- `get()`: Thread-safe. Atomically returns the current value, with memory effects of reading a volatile variable.
 - A context switch can occur right after `get()` returns, and another thread may change the value.
 - However, `AtomicBoolean` guarantees that the returned value of `get()` will be updated and the most up-to-date, whenever it is used later on.

- `public final boolean compareAndSet(boolean expect, boolean update)`

- Atomically sets the “update” (new) value if the current value is equal to the “expect” value.
- Returns true if successful.
- Returns false if the current value was not equal to the “expect” value.

```
- atomicFlag.compareAndSet(true, false); // Thread-safe
  • Sets false if the current value is true, and returns true
  • Keeps the current value if it is false, and returns false
```

AtomicBoolean for Thread Termination

- `AtomicBoolean` can be used to implement a flag in flag-based and 2-step thread termination schemes

```
class CancelableRunnable
implements Runnable {

    boolean done = false;
    ReentrantLock lock;

    public void setDone(){
        lock.lock();
        done = true;
        lock.unlock();
    }

    public void run(){
        while(true){
            lock.lock();
            if(done) break;
            lock.unlock();
            ... // do some work }}}

class CancelableRunnableWithAtomicBoolean
implements Runnable {

    AtomicBoolean done =
        new AtomicBoolean(false);

    public void setDone(){
        done.set(true);
    }

    public void run(){
        boolean temp = false;
        while(true){
            temp = done.get();
            if(done.get()) break;
            ... // do some work
        } } }
```

AtomicInteger

- Offers thread-safe methods to manipulate a single integer value atomically.

```
- int i = 0;
  int j = i;
  if(i==0)                // 2 steps. Not thread-safe.
                        // 3 steps. Not thread-safe.
```

```
- volatile i = 0;
  int j = i;
  if(i==0)                // 2 steps. Thread-safe.
                        // 3 steps. Thread-safe.
```

```
- AtomicInteger atomicInt = new AtomicInteger(0);
  int j = atomicInt.get(); // Thread-safe. j==0
  if(atomicInt.get()==0){...} // Thread-safe.
```

- `get()`: Thread-safe. Atomically returns the current value, with memory effects of reading a volatile variable.

- A context switch can occur right after `get()` returns, and another thread may change the value.
- However, `AtomicInteger` guarantees that the returned value of `get()` will be updated and the most up-to-date, whenever it is used later on.

Exercise (Not HW)

- Define a flag as an `AtomicBoolean` variable in a 2-step thread termination scheme that you have implemented for a prior HW.

```
- volatile i = 0;
  i = 1; // Thread-safe.
  int j = i + 1; // 5 steps. NOT thread-safe.
  i++; // 5 steps. NOT thread-safe.
```

```
- AtomicInteger atomicInt = new AtomicInteger(0);
  atomicInt.set(1); // Thread-safe.
  j = atomicInt.incrementAndGet(); // Thread-safe. j==2
  atomicInt.incrementAndGet(); // Thread-safe.
```

- `set()`: Thread-safe. Atomically returns the current value, with memory effects of reading a volatile variable.

- *read-and-update* methods (e.g., `xxxAndSet()`, `xxxAndGet()` and `getAndXxx()`) have the memory effects of both reading and writing volatile variables.

```

- volatile i = 0;
  i = 1;           // Thread-safe.
  if(i==0){i = 1;} // NOT thread-safe.
                  // Read and write ops are thread-safe
                  // each, but if a context switch occurs
                  // in between the two ops, a race
                  // condition can occur.

- AtomicInteger atomicInt = new AtomicInteger(0);
  atomicInt.compareAndSet(0, 1); // Thread-safe.

```

- Other *read-and-update* methods include:

```

- decrementAndGet(), addAndGet(int), getAndSet(int), ...
- updateAndGet(IntUnaryOperator),
  accumulateAndGet(int, IntBinaryOperator)

```

- **updateAndGet**(IntUnaryOperator updateFunction)
 - *Atomically* updates the current integer value with the result of running a given lambda expression, and then returns the updated value.
 - IntUnaryOperator: a general-purpose functional interface
 - updateFunction: a lambda expression that takes the current int value as a parameter, updates it and returns the updated value.
 - This update logic runs atomically (i.e., in a thread-safe manner)
 - ```

AtomicInteger atomicInt = new AtomicInteger(10);
atomicInt.updateAndGet((int i)->++i); // 11. Thread safe
atomicInt.incrementAndGet(); // 12. Thread safe

```

|                  | Params | Returns | Example use case |
|------------------|--------|---------|------------------|
| UnaryOperator<T> | T      | T       | Logical NOT (!)  |

```

• AtomicInteger atomicInt = new AtomicInteger(10);
 atomicInt.updateAndGet((int i)->++i); // 11. Thread safe

```

- Why ++1? Just in case, note that:

```

- int i = 0;
 i++; // i==1

- int i = 0;
 int x = 0;
 int x = i++; // i==1, x==0

- int i = 0;
 int y = 0;
 int y = ++i; // i==1, y==1

```

```

• volatile i = 10;
 if(i > 0){i = 0;} // NOT thread-safe.
 // Read and write ops are thread-safe
 // each, but if a context switch occurs
 // in between the two ops, a race
 // condition can occur.

• AtomicInteger atomicInt = new AtomicInteger(10);
 atomicInt.updateAndGet((int i)->{ if(i > 0){i = 0;}
 return i; });

// Thread safe. The lambda expression
// is executed in a thread-safe manner.

```

- `accumulateAndGet(int, IntBinaryOperator)`
  - *Atomically* updates the current `int` value with the result of applying a given function, and then returns the updated value.
- `IntBinaryOperator`: a general-purpose functional interface
- `atomicInt.accumulateAndGet(initValue, (result, currentVal)-> ... );`
- `int result = initValue;`  
`result = accumulate(result, currentVal);`
- Takes a lambda expression as the second parameter.
  - The body of `accumulate()` is expressed in the LE, which runs atomically
  - C.f. `Stream.reduce()`

|                                      | Params | Returns | Example use case            |
|--------------------------------------|--------|---------|-----------------------------|
| <code>BinaryOperator&lt;T&gt;</code> | T, T   | T       | Multiplying two numbers (*) |

- `AtomicInteger atomicInt = new AtomicInteger(0);`  
`atomicInt.accumulateAndGet(10, (result, currentVal)-> (currentVal+result)/2 );`
- `AtomicInteger atomicInt = new AtomicInteger(0);`  
`atomicInt.accumulateAndGet(0, (result, currentVal)->{ if(currentVal >= result) return currentVal; else if return result; }));`

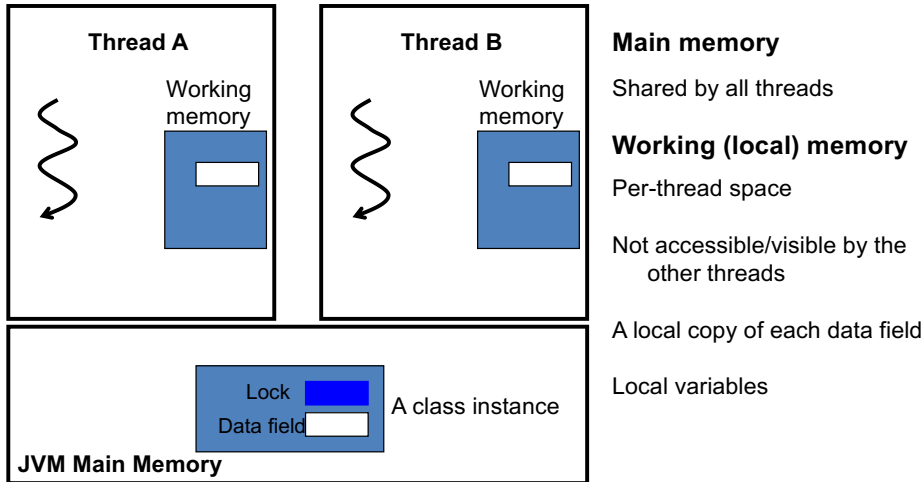
## Appendix: The *Volatile* Keyword and JVM Memory Model

## Memory Management in JVM

- A Java Virtual Machine (JVM) manages memory space for a Java program(s), following [JVM Memory Model](#).



# JVM Memory Model (Java 5~)



33

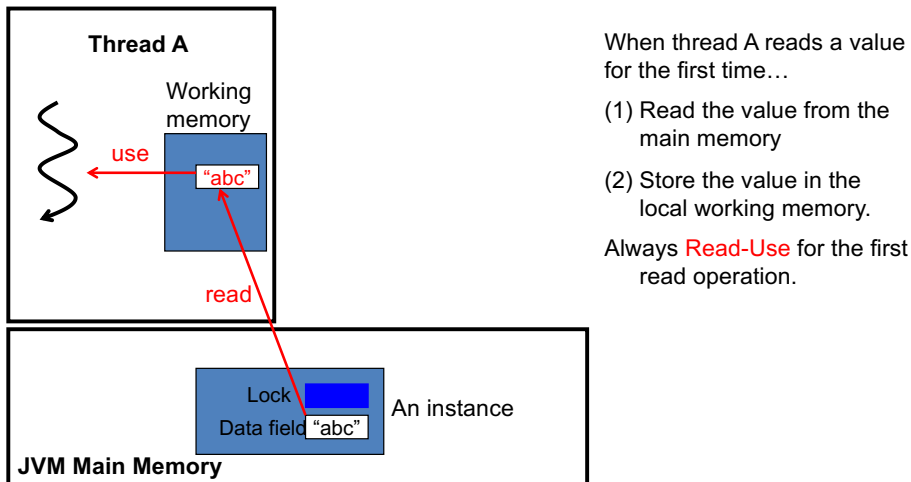
# JVM Actions

- JVM implements a set of *JVM actions* to manage the main memory space and local memory spaces.
- JVM actions (Java bytecode instructions)
  - *read, write, use, assign*
  - *lock, unlock*
  - These are all atomic.

34

## Read Operation (Single Threaded)

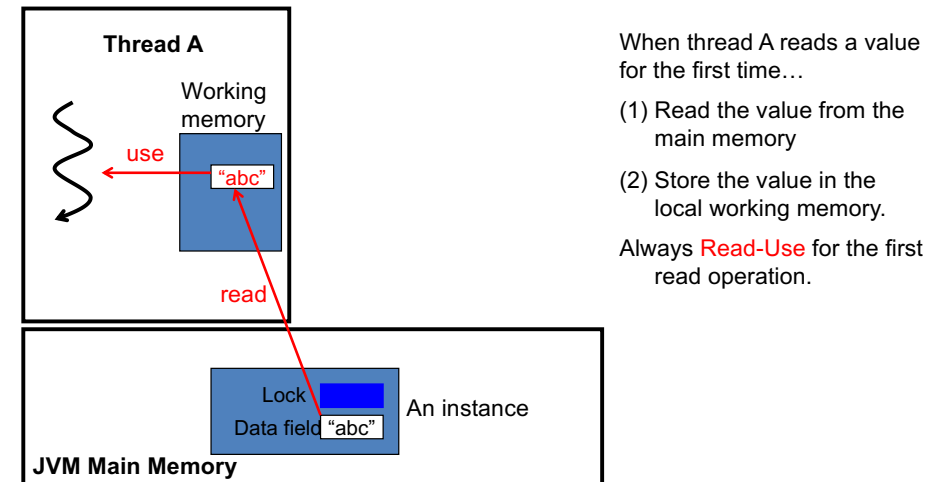
`System.out.println(variable);`



35

## Read Operation (Single Threaded)


`if(variable.equals(...))`



36

## Read Operation (Single Threaded)

```
if(variable.equals(...)){ // First read operation
 System.out.println(variable); } // Second read operation
```



When thread A reads a value for the first time...

- (1) Read the value from the main memory
- (2) Store the value in the local working memory.

Always **Read-Use** for the first read operation.

When thread A accesses the same variable later...


**Use OR Read-Use**

It's up to JVM implementations, but often, **Use** only because it's faster.

37

## Write Operation (Single Threaded)

```
variable = "abc";
```



When thread A assigns a value to a data field...


- (1) Assign the value to a data field in the local memory.
- (2) Copy it to the main memory.

Always **Assign-Write** for the first assignment operation.

38

## Write Operation (Single Threaded)

```
variable = "abc"; // First write operation
variable = "cde"; // Second write operation
```



When thread A assigns a value to a data field...

- (1) Assign the value to a data field in the local memory.
- (2) Copy it to the main mem.

Always **Assign-Write** for the first assignment operation.

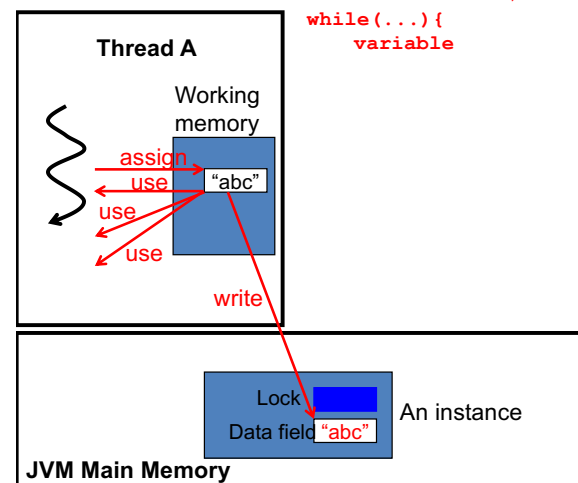
It is not predictable when (how soon) "write" occurs. Maybe immediately after "assign," but some interval may exist.

If thread A repeatedly assigns different values to the same data field, it makes sense to do: **Assign, Assign, ... , Write**

39

## Write Operation (Single Threaded)

```
variable = "abc";
while(...){
 variable
```



When thread A assigns a value to a data field...

- (1) Assign the value to a data field in the local memory.
- (2) Copy it to the main mem.

Always **Assign-Write** for the first assignment operation.

It is not predictable when (how soon) "write" occurs. Maybe immediately after "assign," but some interval may exist.

If thread A repeatedly assigns different values to the same data field, it makes sense to do: **Assign, Assign, ... , Write**

If thread A often read the value, it makes sense to do: **Assign, Use, Use, Use, ... , Write**

40

# Thread Synchronization and Memory Synchronization

- When a thread runs atomic code with a lock held...
  - `public void setDone(){`
    - `lock.lock()`
    - `// atomic code;`
    - `lock.unlock(); }`
- JVM does two things:
  - Thread synchronization
    - Allows only one thread to enter and run atomic code at a time.
    - All the other threads are blocked.
  - Memory synchronization
    - Synchronizes the most up-to-date value in between a local memory and the main memory.

41


## Race Condition

- A race condition can occur due to
  - Failure of thread synchronization and/or
  - Failure of memory synchronization
    - Inconsistency between data in working and main memories.

42

## Race Condition

- A race condition can occur due to
  - Failure of thread synchronization and/or
  - Failure of memory synchronization
    - Inconsistency between data in working and main memories.



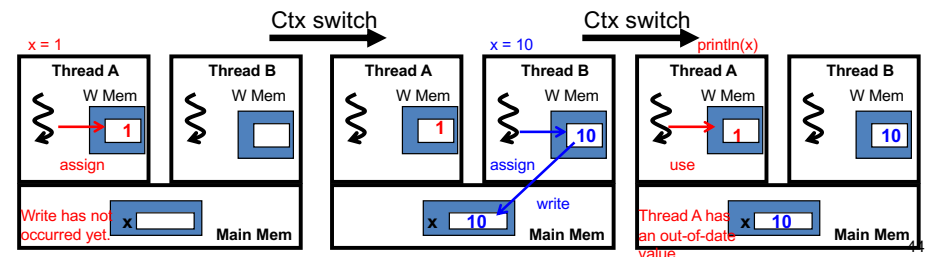
*A race condition can occur due to a failure of thread synchronization.*

43

## Race Condition

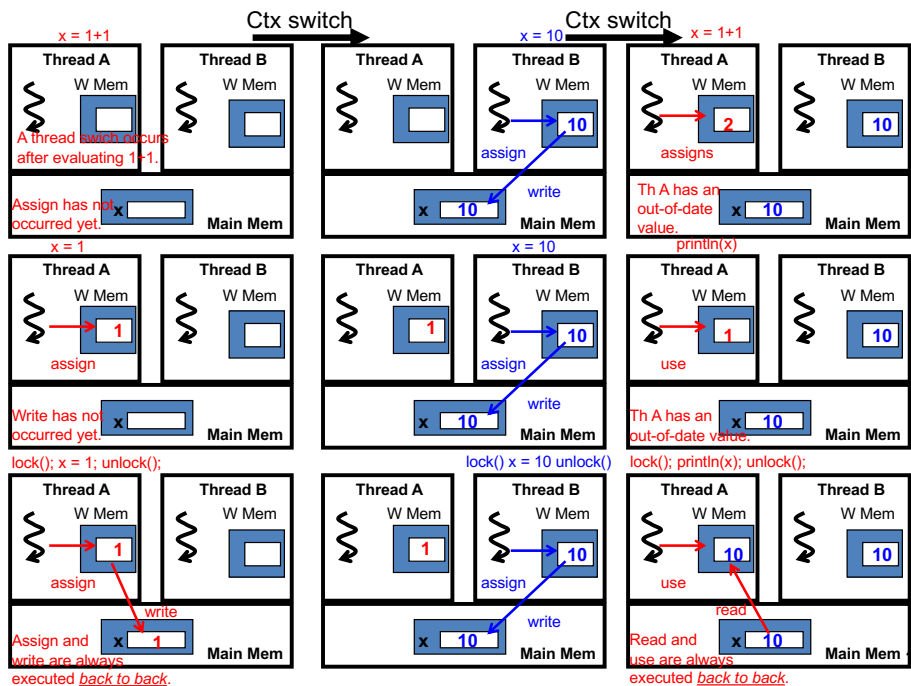
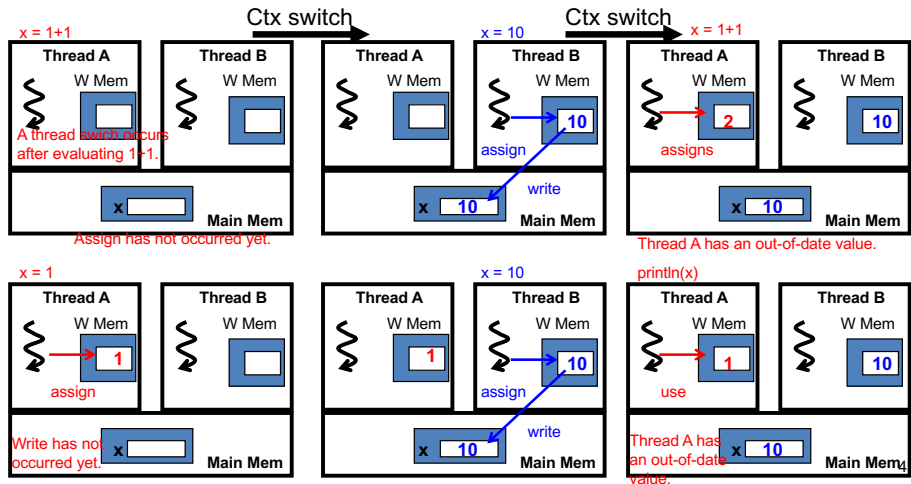
- A race condition can occur due to
  - Failure of thread synchronization and/or
  - Failure of memory synchronization
    - Inconsistency between data in working and main memories.

*A race condition can occur due to a failure of memory synchronization. Threads are synchronized in this case (by chance).*



## Race Condition

- A race condition can occur due to
  - Failure of thread synchronization
  - Failure of memory synchronization
    - Inconsistency between data in working and main memories.



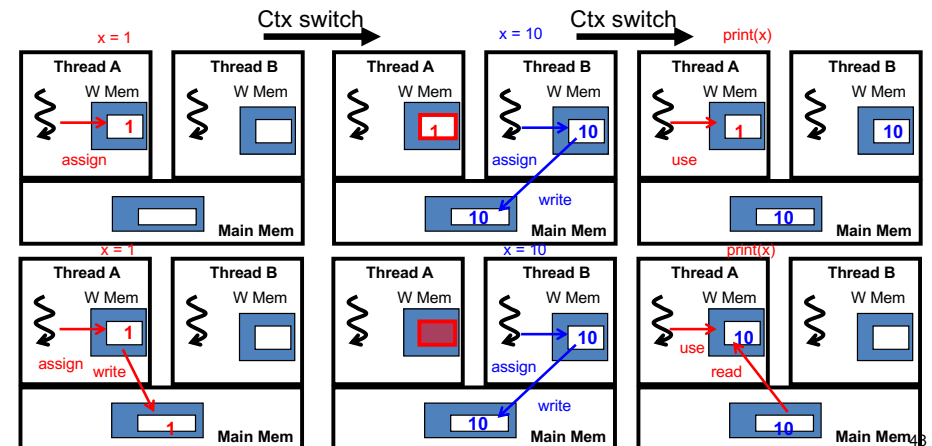
## Race Conditions and ReentrantLock

- You need **both thread and memory synchronization** to prevent race conditions.
- ReentrantLock** does both.
  - Thread synchronization
    - Allows only one thread to enter and run atomic code at a time.
      - All the other threads are blocked.
  - Memory synchronization
    - Synchronizes the most up-to-date value in between a local memory and the main memory.
      - Destroys working memory upon entering atomic code
      - Flushes working memory to the main memory upon exiting atomic code

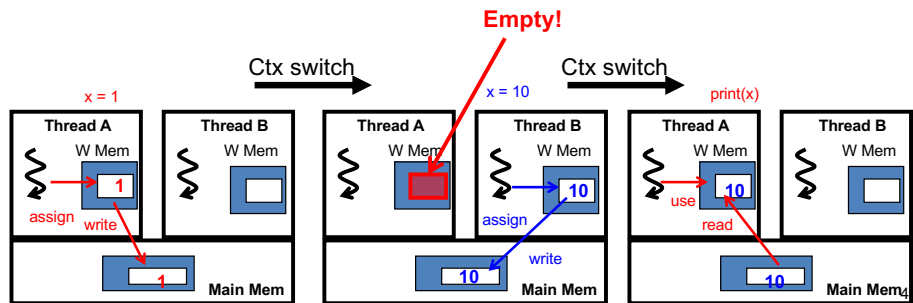
46

## Volatile Variables

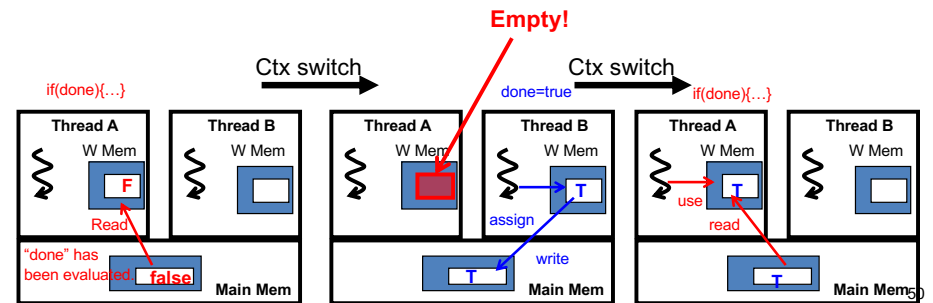
- When a thread uses a volatile variable...
  - Memory synchronization is guaranteed.



- Write always follows Assign immediately.
  - They are always paired.
- Read always occur before Use.
  - They are always paired.
- A volatile variable's value is *NOT persistent* (i.e., *volatile*) across context switches.



- Write always follows Assign immediately.
  - They are always paired.
- Read always occur before Use.
  - They are always paired.
- A volatile variable's value is *NOT persistent* (i.e., *volatile*) across context switches.



## When to use Volatile Variables?

- No need to use locking to perform memory synchronization for a volatile variable.
  - No locking → less overhead
- A volatile variable NEVER synchronizes threads that access the variable.
  - “Volatile” is not a thread sync tool. It’s for a memory sync tool.
- ***The volatile keyword works when you don’t need to synchronize threads but need to synchronize memory.***
- Latch
  - A data structure that performs a single type of state changes
    - e.g. False → True
  - Often used to terminate threads.
    - c.f. “done” variable in prior examples

# An Example Latch

- `volatile` boolean done = false;
- ```
public void setDone(){
    done = true;}
```
- ```
public void run(){
 while(true)
 if(done) break;
 counter++;
}
```
- The state of “done” always changes in a unidirectional way:  
false → true
  - “true → false” never occur.

53

- `volatile` boolean done = false;
- ```
public void setDone(){
    done = true;}
```
- ```
public void run(){
 while(true)
 if(done) break;
 counter++; }
```
- No need to surround the if statement with lock() and unlock().
- Thread sync is not performed.
  - A context switch can occur in between
    - Evaluating the “done” variable and
    - Applying the current value in “done” to the if statement
- Memory sync is performed.
  - The most up-to-date value of “done” is applied to the if statement.

54

## Syntactic Difference

- `volatile` boolean done = false;
- ```
public void setDone(){
    done = true;} // NO NEED TO SURROUND THIS WITH LOCK() and UNLOCK()
```
- ```
public void run(){
 while(true)
 if(done) break;
 counter++; }
```
- Thread sync is not performed.
  - A context switch can occur in between evaluating the value of “true” and assigning it “done.”
  - All threads will assign “true” to “done.”
    - No other possible state changes.
  - Writer threads do not generate race conditions.
- Memory sync is performed.
  - The value of “true” must be copied to the main memory once the assignment (“done=true”) is completed.

55

- | • Without “volatile”                                                                                                                                            | • With “volatile”                                                                                                           |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------|
| <pre>boolean done = false; ReentrantLock lock =     new ReentrantLock();  public void setDone(){     lock.lock();     done = true;     lock.unlock(); }</pre>   | <pre><code>volatile</code> boolean <code>done</code> = false;  public void setDone(){     <code>done</code> = true; }</pre> |
| <pre>public void run(){     while( true ){         lock.lock();         if( <code>done</code> ) break;         // Do some work         lock.unlock(); } }</pre> | <pre>public void run(){     while( true )         if( <code>done</code> ) break;         // Do some work }</pre>            |

56

# Limited Effectiveness/Usefulness

- A “volatile” variable is **guaranteed to have the most up-to-date value** whenever it is read.

```
- volatile int a;
int b = a; // These 2-step operations are all thread-safe,
if(a==0){...} // even if a context switch occurs in between
println(a); // the 2 steps and another thread changes the
 // value of "a" there.
```

- However, it is not a silver bullet to eliminate all possible race conditions.

```
- a + 1; // 3 steps. Thread-safe.
b = a + 1; // These 4-step operations are NOT thread-safe.
if(a+1>0){...}; // The first 3 steps are thread-safe though. A
 // race condition can occur in between the 3rd
 // and 4th steps.

println(a+1) // 4 steps. Not thread-safe.
```

- “volatile” is effective only in the read operations that have **no intermediate state**.

57

- A “volatile” variable **NEVER** eliminate race conditions in write operations on it.

```
- volatile int a;
a = 1; // 2 steps. A race condition can occur if a
 // context switch occurs in between the 2 steps
 // and another thread changes the value of "a"

a = a + 1; // 4 steps. NOT thread-safe
a++;
```

- Use a “volatile” variable only when you can live with these race conditions.
- Do NOT use “volatile” for arrays.

## A Concurrency Bug in Jetty

- Jetty
  - An open source web implementation in Java
    - <http://jetty.codehaus.org/jetty/>
- A bug report (March 2010)

```
- // Jetty 7.1.0,
// org.eclipse.jetty.io.nio,
// SelectorManager.java, line 105
private volatile int _set;
public void register(SocketChannel channel, Object att){
 int s = _set++; // This part is NOT thread-safe
 ...
}
public void addChange(Object point){
 synchronized (_changes){
 ...
 }
}
```

## In Summary...

- NOT a general-purpose, widely-applicable threading tool
- Powerful only in some specific cases
  - In practice, assume it is useful only for simplifying the implementation of a latch.
    - Useful to implement flag-based thread termination and 2-step thread termination.