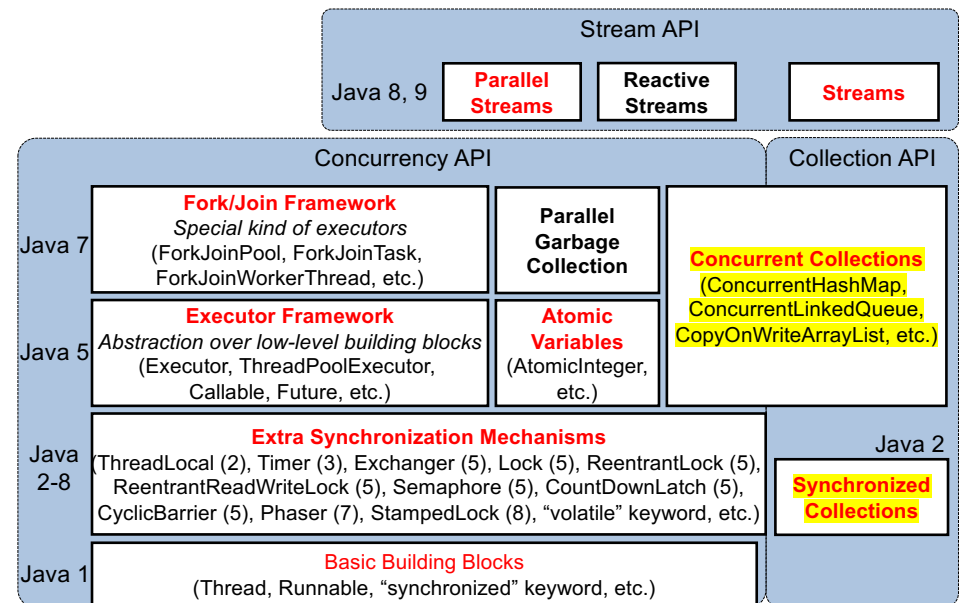


## Collections and Concurrency

## Concurrency API in Java



## 3 Types of Collections in Java

- Thread-unsafe collections
- Thread-safe collections
  - Synchronized collections
  - Concurrent collections

## Thread-unsafe Collections

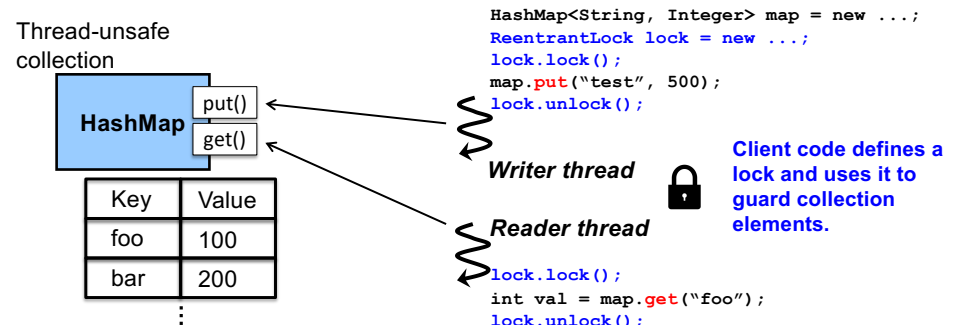
- Many collection classes are **NOT thread safe**.
  - e.g., ArrayList, LinkedList, HashMap, etc.
  - Their public methods **never** perform thread synchronization (locking).
- Look into Java API documentation to see if a collection is thread-safe or not.

# Java API Doc on ArrayList

- “Note that this implementation is **not synchronized**. If multiple threads access an ArrayList instance concurrently, and at least one of the threads modifies the list **structurally**, it *must* be synchronized **externally**.
  - (A structural modification is any operation that adds or deletes one or more elements, or explicitly resizes the backing array; merely setting the value of an element is not a structural modification.)...”

# When You Use a Thread-unsafe Collection...

- You must do **client-side locking**; your client code must **externally** perform **thread synchronization**.
  - to guard collection elements against concurrent accesses.
    - c.f. previous HWs that use thread-unsafe collections such as ArrayList, LinkedList and HashMap.



# Example Compound Operations

- You must do **client-side locking** for **compound operations** as well as simple public method calls.
- Example compound operations
  - Iteration (element-traversal)
    - Get all elements one by one
  - Navigation
    - Find/search the next element after a given element
  - Conditional operations (check-then-act)
    - e.g., Check if a **HashMap** has a key-value pair for the key **k**, and if not, add the pair **(k,v)**

- ```
List<String> list = new ArrayList<String>();
...

Iterator it = list.iterator();
while( it.hasNext() ){
    doSomething( it.next() ); } // Iteration

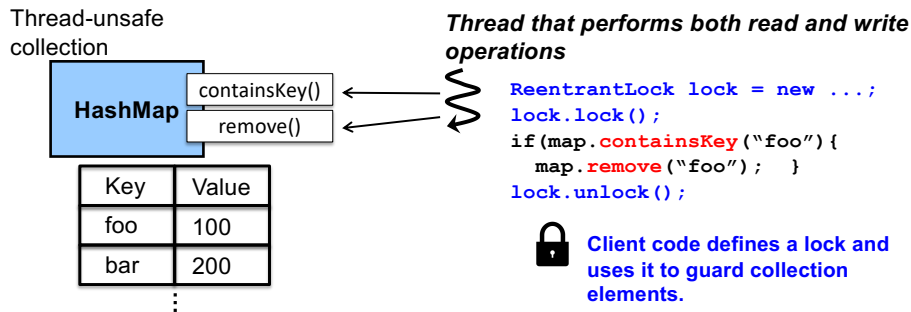
for(int i=0; i<list.size(); i+=2){ // Navigation
    doSomething( list.get(i) ); }

while( it.hasNext() ){
    if( it.next().endsWith(".jar" ) {... } // Navigation

if( list.size() > 10 ) // Check-then-act
    doSomething( list );
```

# Client-side Locking for Compound Operations

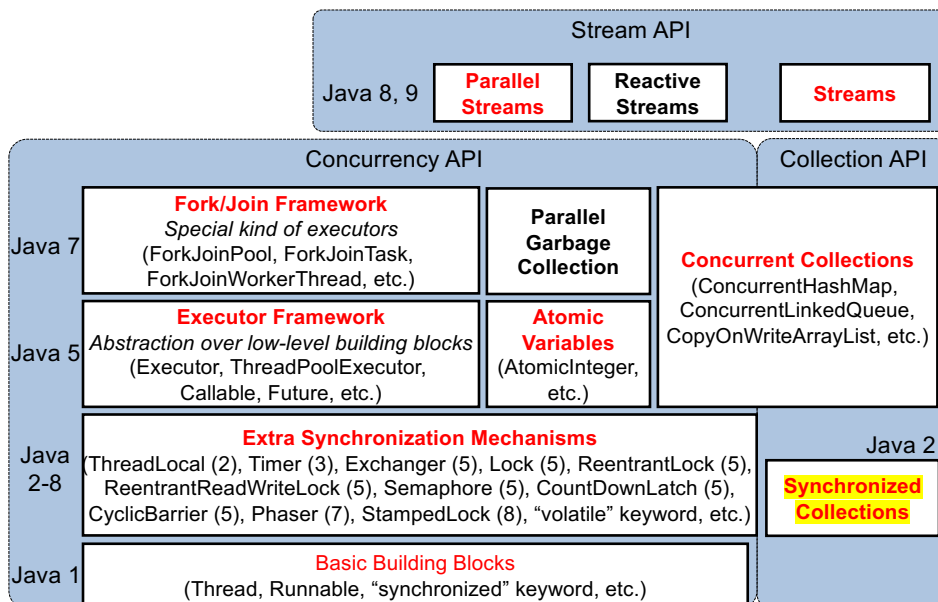
- An example “check-then-act” compound operation



# 3 Types of Collections in Java

- Thread-unsafe collections
- Thread-safe collections
  - Synchronized collections
  - Concurrent collections

## Concurrency API in Java



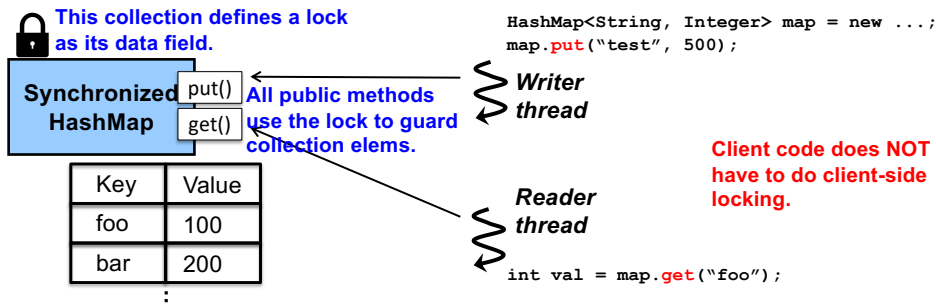
## Synchronized Collections

- “Ready-made” thread-safe collections
  - Synchronized classes: Vector and Hashtable
    - All public methods perform thread synchronization.
      - Only one thread can access the elements of a collection at a time.
        - » e.g., When a thread is in the middle of executing add() on a Vector, no other threads can call get(), size(), etc. on that Vector.
  - Synchronized wrapper classes for thread-unsafe collections
    - Created by java.util.Collections.synchronizedXyz()
      - Factory methods
      - synchronizedList()
      - synchronizedMap()
      - synchronizedSet()

## Vector and Hashtable in Single Threaded Programs

- It makes no sense to use these collections in single-threaded programs in a **performance point of view**.
  - They perform thread synchronization even when only one thread runs in a program.
    - Unnecessary performance loss
- Use **ArrayList** and **HashMap** instead!
  - They still exist in Java API only for backward compatibility.

- All public methods are **thread-safe** in all synchronized collection classes.
- No client-side locking is necessary when client code makes a simple public method call.



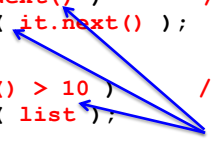
## Synchronized Wrapper Classes

- `List<String> list = Collections.synchronizedList( new ArrayList<String>() );`
- `list`: an instance of a synchronized wrapper class for `ArrayList`.
  - `list.getClass()` → `java.util.Collections$SynchronizedRandomAccessList`
  - The wrapper class offers “synchronized” (or **thread-safe**) versions of `ArrayList`’s public methods.
    - `add()`, `get()`, `remove()`, etc.

- Need **client-side locking in compound operations** on a synchronized collection
  - **Iteration (element-traversal)**
    - Get all elements one by one
  - **Navigation**
    - Find/search the next element after a given element
  - **Conditional operations (check-then-act)**
    - e.g., Check if a **HashMap** has a key-value pair for the key `k`, and if not, add the pair `(k, v)`

## Potential Problems in Compound Actions

```
• List<String> list =  
  Collections.synchronizedList( new ArrayList<String>());  
  
  Iterator it = list.iterator();  
  while( it.hasNext() )      // Iteration  
    doSomething( it.next() );  
  
• if( list.size() > 10 )      // Check-then-act  
  doSomething( list );
```



Race conditions can occur here.

- Race conditions
- ConcurrentModificationException
  - Raised if a writer thread tries to add/remove elements before a reader thread completes a traversal on the entire set of elements.

```
• ReentrantLock lock = new ReentrantLock();  
  lock.lock();                // acquires lock  
  Iterator it = list.iterator();  
  while( it.hasNext() )      // acquires the lock that the list owns  
    doSomething(it.next()); // acquires the lock that the list owns  
  lock.unlock();
```

- `lock` is different from the lock that the `list` owns for thread synch in its public methods.
  - Must make sure to use `lock` consistently in all client code of `list`.

- A thread acquires 2 locks.

## Client-side Locking for Compound Operations

```
• List<String> list =  
  Collections.synchronizedList( new ArrayList<String>());
```

```
• synchronized(list){  
  Iterator it = list.iterator();  
  while( it.hasNext() )      // nested locking  
    doSomething( it.next() ); // nested locking  
}
```

```
• synchronized(list){  
  if( list.size() > 10 )      // nested locking  
    doSomething( ... );  
}
```

- `synchronized(list)`: acquires the lock that the `list` owns/uses for thread synchronization in its public methods.

## Performance Implication on Client-side Locking

```
• List<String> list =  
  Collections.synchronizedList( new ArrayList<String>());
```

```
• synchronized(list){          // acquires the lock that the list owns  
  Iterator it = list.iterator();  
  while( it.hasNext() )      // nested locking  
    doSomething( it.next() ); // nested locking  
}
```

```
• ReentrantLock lock = new ReentrantLock();  
  lock.lock();                // acquires lock  
  Iterator it = list.iterator();  
  while( it.hasNext() )      // acquires the lock that the list owns  
    doSomething( it.next() ); // acquires the lock that the list owns  
  lock.unlock();
```

- Performance penalty due to 2 lock acquisitions

## Summary: Thread-unsafe Collections v.s. Synchronized Collections

- `List<String> list = new ArrayList<String>();`  
`...`  
`ReentrantLock lock = new ReentrantLock();`  
`lock.lock();`  
`Iterator it = list.iterator();`  
`while( it.hasNext() )`  
`doSomething( it.next() );`  
`lock.unlock();`
- Performs client-side locking for **ArrayList**, which is **NOT** thread-safe.
  - The client code is thread-safe although `hasNext()` and `next()` are not.
- **More efficient** than the previous client code.
  - Only 1 lock acquisition (not 2 acquisitions)
  - Use a read-write lock if you have many “reader” threads.

### Thread unsafe collections

- Client code always must perform thread synchronization.
  - For **both simple public method calls and compound operations**
  - Client-side locking always requires 1 lock acquisition.

### Synchronized collections

- Client code does NOT have to perform thread synchronization **for simple public method calls**.
- However, it needs client-side locking **for compound operations**.
  - Client-side locking requires 2 lock acquisitions.

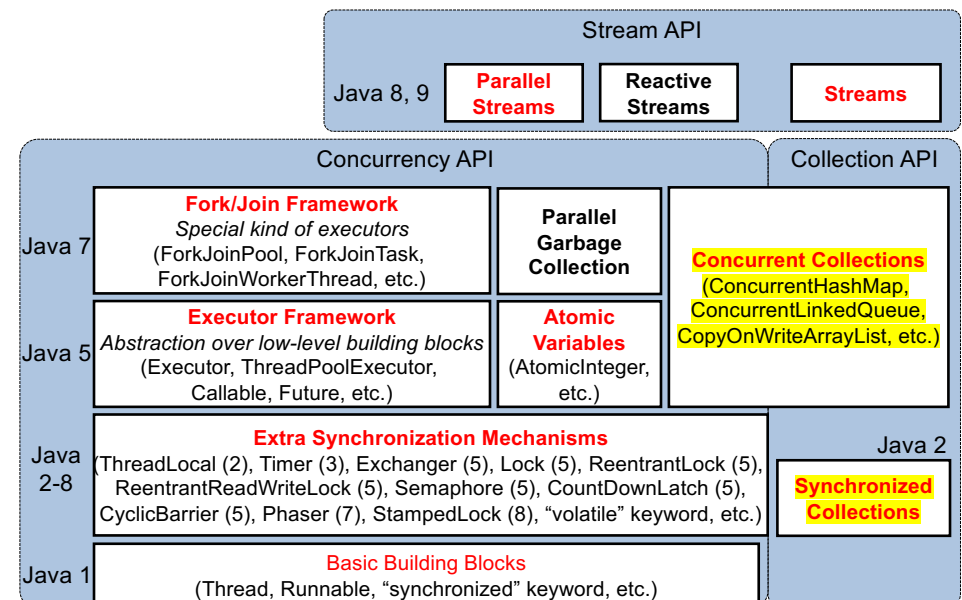
### Thread unsafe collections

- Use them consistently.

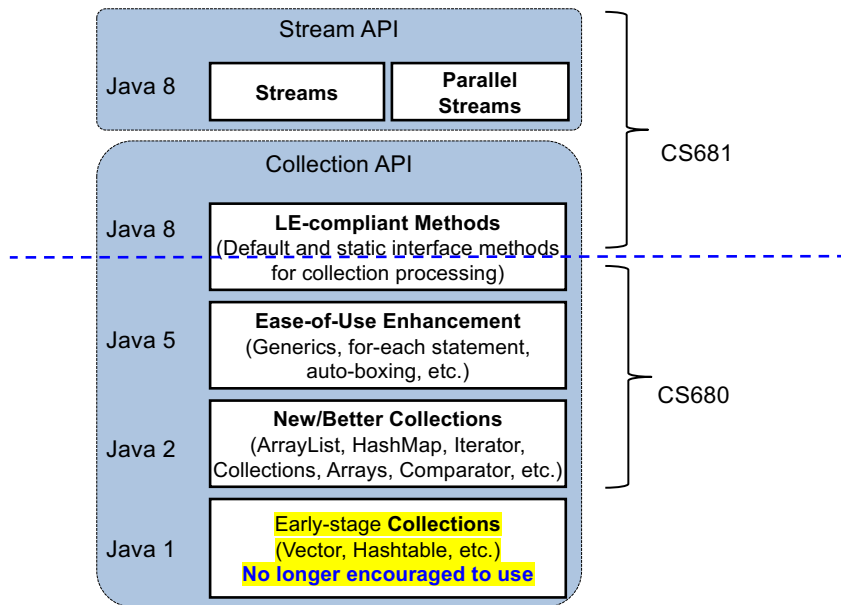
### Synchronized collections

- Forget about them.
- Never use `Vector` and `Hashtable` in single-threaded programs

## Concurrency API in Java



## Collection API in Java

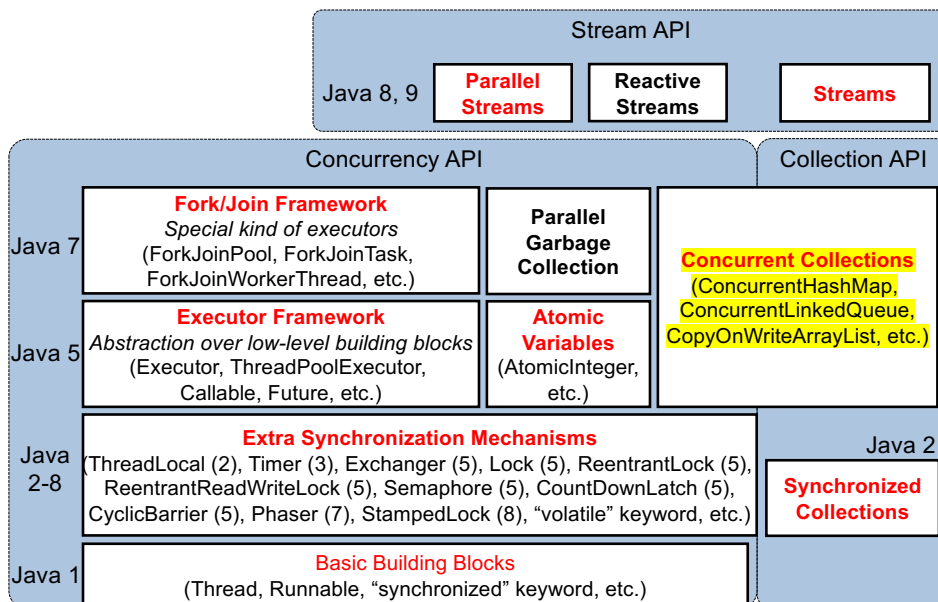


25

## 3 Types of Collections in Java

- **Thread-unsafe** collections
- **Thread-safe** collections
  - **Synchronized** collections
  - **Concurrent** collections

## Concurrency API in Java



## Concurrent Collections

- “Ready-made” **thread-safe** collections
  - Introduced in Java 5 (2004) and enhanced in subsequent versions
    - **Queue**
      - **ConcurrentLinkedQueue** (since Java 5)
      - **ConcurrentLinkedDeque** (since Java 7)
      - **ArrayBlockingQueue** (since Java 5)
      - **LinkedBlockingQueue** (since Java 5)
      - **DelayQueue** (since Java 5)
      - **PriorirtyBlockingQueue** (since Java 5)
      - **LinkedBlockingDeque** (since Java 6)
      - **LinkedTransferQueue** (since Java 7)
    - **Map**
      - **ConcurrentHashMap** (since Java 5)
      - **ConcurrentSkipListMap** (since Java 6)
    - **Set**
      - **ConcurrentSkipListSet** (since Java 6)
  - **java.util.concurrent.CopyOnWriteXyz** classes
    - **CopyOnWriteArrayList**
    - **CopyOnWriteArraySet**

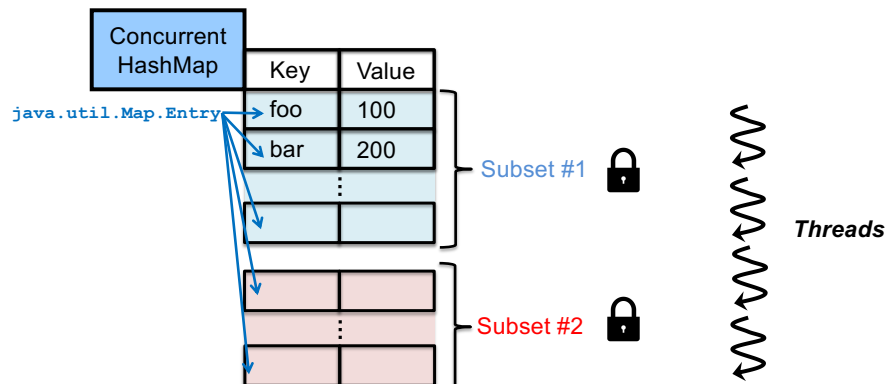
# ConcurrentHashMap

- Indented to **replace synchronized collections**.
  - e.g., `ConcurrentHashMap` is indented to replace `SynchronizedMap`.
- Make public methods **thread-safe**.
  - e.g., `get()` and `put()` in `ConcurrentHashMap`
  - Client code does NOT have to perform thread synchronization **for simple public method calls**.
- Implement public methods in an **efficient** manner.
  - e.g., lock stripping
- Provide public **thread-safe methods** to perform **compound operations**
  - Client code does not have to do client-side locking.

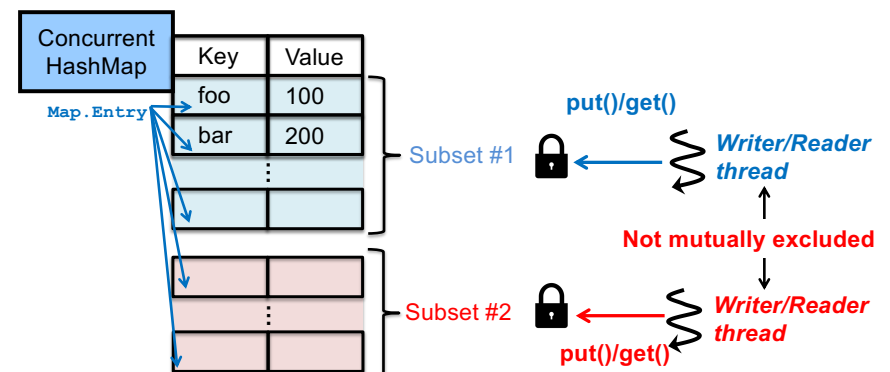
- Provides thread-safe public methods.
  - e.g., `get()` and `put()`
  - Client code does NOT have to do client-side locking.
- A replacement for synchronized hash-based `Map` implementations (e.g., `Hashtable` and synchronized `HashMap`).
- Implements public methods in an **efficient** manner.
  - Performs ***fine-grained locking***, called **lock stripping**
    - Compared to *coarse-grained locking* (i.e., lock-per-collection) in synchronized hash-based `Map` implementations.

## Lock Stripping

- `ConcurrentHashMap` uses **multiple** locks to guard a table (i.e., key-value pairs).
  - 16 locks by default
    - Configurable with the “concurrencyLevel” parameter in a constructor.
  - Each lock is associated with a subset of the table.

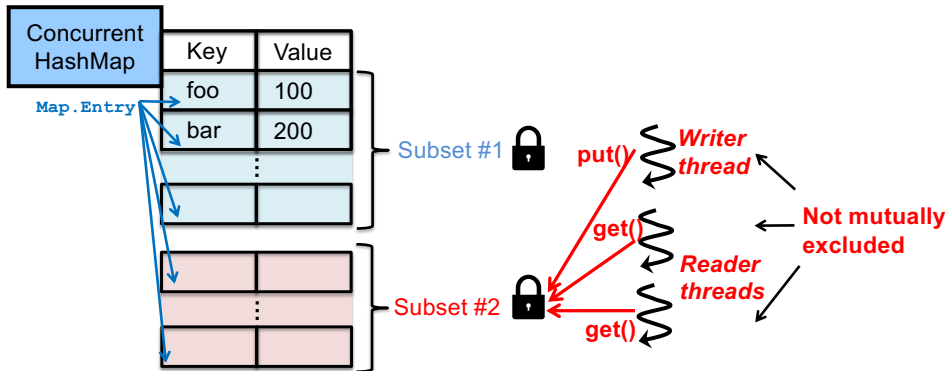


- Threads are **not mutually excluded** with each other
  - as far as they access different subsets of the table with different locks.

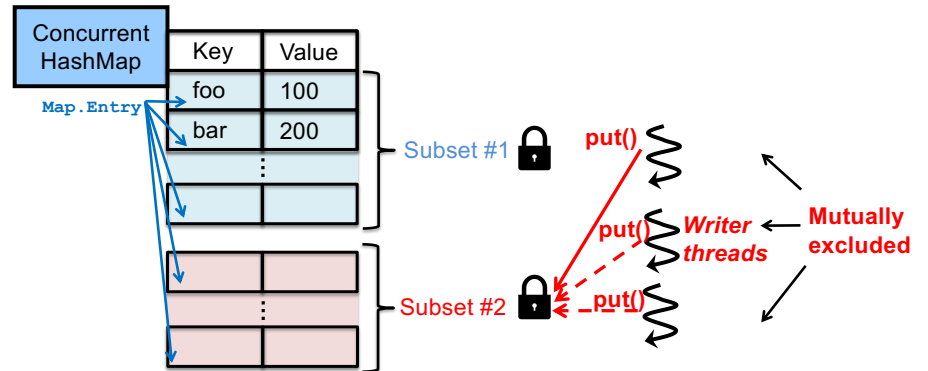




- To access a subset of the table,
  - Reader threads
    - are **NOT mutually excluded** with each other.
      - c.f. read-write lock
    - are **NOT mutually excluded** with writer threads.
      - c.f. inner class Node<K, V>

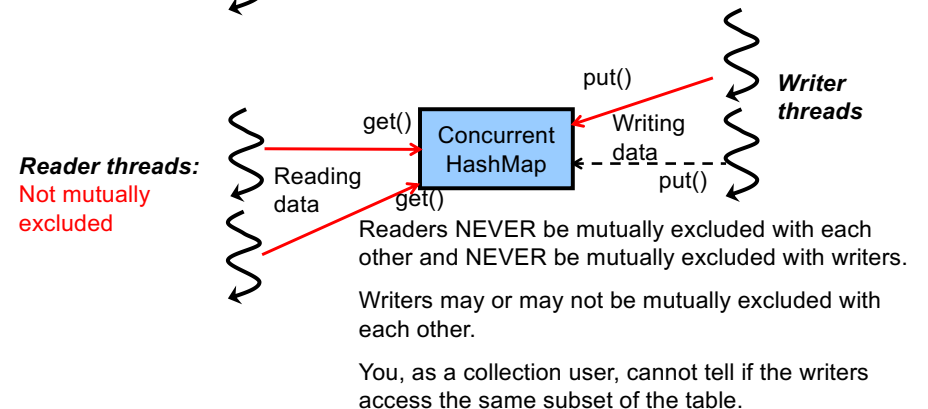
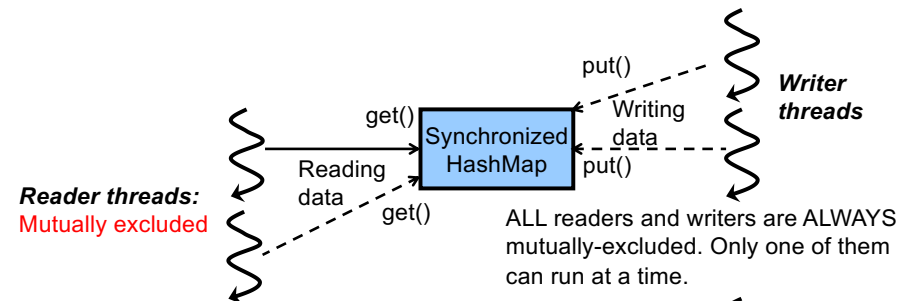



- To access a subset of the table,
  - Writer threads **ARE mutually excluded** with each other.



## Synchronized Hash-based Map Impls

- A **single lock** is used to guard the **entire** table in `Hashtable` and `synchronized HashMap`
  - No lock stripping.
- All writer/reader threads **ARE ALWAYS mutually excluded** with each other.
  - A potential performance bottleneck as the number of key-value pairs increases and the number of threads increases.



# Thread-safe Compound Operations


- Supports **thread-safe iteration**
  - **No client-side locking** is necessary.
    - c.f. Thread-unsafe collections and synchronized collections require client-side locking for iteration
      - Because it is a compound operation.

```

• Iterator it = aConcurrentHashMap.entrySet().iterator();
while( it.hasNext() )
    doSomething( it.next() );
  
```

## – Pros

- **No client-side locking** is necessary.
  - No need to mutually exclude readers.



# Concurrent Iterators

- **Concurrent iterators** are obtained through `entrySet()`, `keySet()` and `values()`.

- `entrySet()`: Returns key-value pairs as a Set.
  - `Set<Map.Entry<K,V>>`
- `keySet()`: Returns keys as a Set.
  - `ConcurrentHashMap.KeySetView<K,V>`
- `values()`: Returns values as a Collection.
  - `Collection<V>`

```


- Iterator it = aConcurrentHashMap.entrySet().iterator();
  while( it.hasNext() )
    doSomething( it.next() );
  
```

```

• Iterator it = aConcurrentHashMap.entrySet().iterator();
while( it.hasNext() )
    doSomething( it.next() );
  
```

## – Pros

- **No client-side locking** is necessary in client code.
  - No need to mutually exclude readers.
  - No need to mutually exclude readers and writers.
    - » Writers can add/remove elements while readers read elements.
    - » It is guaranteed that writers and readers do not corrupt elements.
- The iterator “it” is **backed by the map**, so changes to the map are reflected in the set.




```

• Iterator it = aConcurrentHashMap.entrySet().iterator();
while( it.hasNext() )
    doSomething( it.next() );

```

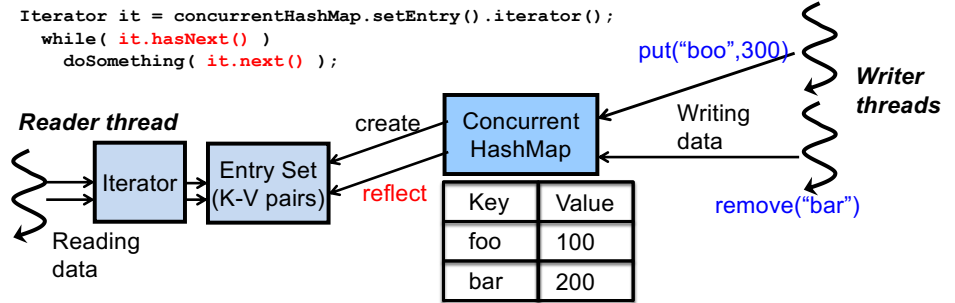
## – Cons:

- *Weak (or best-effort) consistency*
  - There is no guarantee about how soon a change (on the hash map) to be reflected into the entry set.
  - The iterator “it” may or may not traverse the most up-to-date key-value pairs in the map.



## Consistency v.s. Performance

- **ConcurrentHashMap** trades *perfect consistency* for *performance improvement*.
  - If you can live with weak consistency, it's a great collection class for you.
    - Pros: performance improvement
    - Cons:
      - » Iterators may or may not traverse the most up-to-date key-value pairs in the map.
      - » `mappingCount()` and `isEmpty()` are not perfectly reliable.
        - The value returned is an estimate; the actual value may differ if there are concurrent insertions or removals.
  - If you cannot, craft your own thread-safe hash map with **HashMap** and **ReentrantLock**.
    - **ConcurrentHashMap** does not implement perfect consistency.



- The iterator “it” may traverse
  - {(foo,100), (bar,200)},
  - {(foo,100), (bar,200), (boo,300)},
  - {(foo,100)}, or
  - {(foo,100), (boo,300)}.
- It is guaranteed that elements are never get corrupted.
  - Corrupted data like (foo,300) is never put to the map.

- **int size()**
  - Returns the total number of key-value pairs with `int`.
    - int: 32-bit signed integer:- 2147483647 to 2,147,483,647
    - What if you need to have more than 2.15 billion pairs?
- **long mappingCount()** [since Java 8]
  - Returns the total number of key-value pairs with `long`.
    - Long: 64-bit signed integer: -9223372036854775808 to 9,223,372,036,854,775,807
  - Use this method, rather than `size()`, when you maintain a huge number of key-value pairs.

# A New Method in Iterator

- `java.util.Iterator<E>`
  - Used to traverse individual elements in a collection
    - `Iterator it = concurrentHashMap.setEntry().iterator();`  
`while( it.hasNext() )`  
`doSomething( it.next() );`
- `forEachRemaining()` [since Java 8]
  - Accepts a lambda expression (LE) and applies the LE to each element
    - until all elements have been processed or the action throws an exception.
  - `Iterator it = concurrentHashMap.setEntry().iterator();`  
`it.forEachRemaining( (Map.Entry<String, AtomicInteger> e)`  
`->{System.out.println(e); } );`
  - `Iterator it = concurrentHashMap.keySet().iterator();`  
`it.forEachRemaining( (String k)`  
`->{System.out.println(k); } );`
  - No client-side locking is required to call `forEachRemaining()`

# Other Thread-safe Compound Operations

- Supports *common compound operations* in a thread-safe manner
  - put-if-absent: `putIfAbsent(key, value)`
    - Insert a pair of `key` and `value` as a new entry if `key` is not already associated with a value.
  - Conditional remove: `remove(key, value)`
    - Remove the entry for `key` if `key` is associated with `value`.
  - Conditional replace: `replace(key, value)`
    - Replace the entry for `key` if `key` is associated with some value.
  - Conditional replace: `replace(key, oldValue, newValue)`
    - Replace the entry for `key` with `newValue` only if `key` is associated with `oldValue`.
  - No client-side locking is necessary


# A Specialized Iterator: Spliterator

- `public interface Iterable<T>{`  
`Iterator<T> iterator();`  
`Spliterator<T> spliterator();`  
`... }`
- `java.util.Spliterator<E>`
  - “Split” + “iterator”
  - Can *split* (or *subset*) the entire set of elements and traverse a subset of the elements.
    - `Iterator it = concurrentHashMap.keySet().iterator();`  
`it.forEachRemaining( (String k)`  
`->{System.out.println(k); } );`
    - `Spliterator st1 = concurrentHashMap.keySet().spliterator();`  
`Spliterator st2 = st1.trySplit();`  
`st1.forEachRemaining( (String k)`  
`->{System.out.println(k); } );`  
`st2.forEachRemaining( (String k)`  
`->{System.out.println(k); } );`
    - `st1` and `st2` cover two different subsets.

# Exercise: Concurrent Access Counter

- **AccessCounter**
  - c.f. HW 11 (w/ `HashMap` and `ReentrantLock`)
  - c.f. HW 13 (w/ `HashMap` and `ReentrantReadWriteLock`)
  - Maintains a map that pairs a relative file path and its access count.
  - `increment()`
    - accepts a file path and increments the file’s access count.
  - `getCount()`
    - accepts a file path and returns the file’s access count.

Multi-threaded server



## Thread-safe Methods that Accept Lambda Expressions

- Revise the access counter with `concurrentHashMap`

```
- increment()
  • Integer oldCount = concurrentMap.get(aFilePath);
    int newCount = oldCount==null? 1: oldCount+1;
    concurrentMap.put(aFileName, newCount);
```

- This client code is NOT thread-safe.
  - because it performs a compound operation.
  - although `get()` and `put()` are thread-safe.
  - Need **client-side locking** with a `ReentrantLock`

- Alternatively, use `putIfAbsent()` and `AtomicInteger`

```
- ConcurrentHashMap<Path, AtomicInteger> map = new ...;
  map.putIfAbsent(aFilePath, new AtomicInteger(0));
  map.get(aFileName).incrementAndGet();
```

- `compute(K, BiFunction<K,V,V>)`
  - Accepts a key ( $\kappa$ ) and a lambda expression (LE)
  - Applies the LE on a key-value pair that is associated with the supplied key ( $\kappa$ ).
- `map.compute(aFilePath, (Path k, Integer v)->{return v==null? 1: ++v;})`
  - If a key-value pair that is associated with  $\kappa$  does not exist, `null` is passed to the 2nd parameter of the lambda expression.
- `computeIfAbsent(K, Function<K,V>)`
  - `map.computeIfAbsent(aFilePath, (Path k)->{return 1;})`
- `computeIfPresent(K, BiFunction<K,V,V>)`
  - `map.computeIfPresent(aFilePath, (Path k, int v)->{return ++v;})`

## Important General-Purpose Functional Interfaces

|                                      | Params | Returns | Example use case                                                                    |
|--------------------------------------|--------|---------|-------------------------------------------------------------------------------------|
| <code>Function&lt;T,R&gt;</code>     | T      | R       | Get the price (R) from a Car object (T)<br>Generate a function (R) from another (T) |
| <code>Consumer&lt;T&gt;</code>       | T      | void    | Print out a collection element (T)                                                  |
| <code>Predicate&lt;T&gt;</code>      | T      | boolean | Has this car (T) had an accident?                                                   |
| <code>Supplier&lt;T&gt;</code>       | NO     | T       | A factory method. Create a Car object and return it.                                |
| <code>UnaryOperator&lt;T&gt;</code>  | T      | T       | Logical NOT (!)                                                                     |
| <code>BinaryOperator&lt;T&gt;</code> | T, T   | T       | Multiplying two numbers (*)                                                         |
| <code>BiFunction&lt;U,T&gt;</code>   | U, T   | R       | Return TRUE (R) if two params (U and T) match.                                      |

## HW 17

- Revise your concurrent access counter (HW 11 solution) with `concurrentHashMap`
  - Eliminate client-locking (to guard a hash map) in `increment()` and `getCount()`
    - i.e., `AccessCounter` no longer needs a `ReentrantLock` data field.
  - Use lambda expressions whenever possible.

# Bulk Operation Methods

- Apply a given lambda expression on each key-value pair **with multiple threads** in a thread-safe manner.

– **forEach**: `forEach()`, `forEachEntry()`, `forEachKey()`, `forEachValue()`

– **Search**: `search()`, `searchEntries()`, `searchKeys()`, `searchValues()`

– **Reduce**:

- `reduce()`, `reduceToDouble()`, `reduceToInt()`, `reduceToLong()`
- `reduceEntries()`, `reduceEntriesToDouble()`, `reduceEntriesToInt()`, `reduceEntriesToLong()`
- `reduceKeys()`, `reduceKeysToDouble()`, `reduceKeysToInt()`, `reduceKeysToLong()`
- `reduceValues()`, `reduceValuesToDouble()`, `reduceValuesToInt()`, `reduceValuesToLong()`

## forEach Bulk Operations

- Performs a given lambda expression (LE) on each key-value pair.

– `map.forEach(500, (k,v)->{if(v>10000) System.out.println(k + "->" + v) })`

- 1st param: parallelism threshold (500)
- 2nd param: lambda expression (`BiConsumer`)
  - Printing out the key-value pairs that have more than 10,000 access count.

- Each bulk operation method receives a “**parallelism threshold**” as the first parameter and creates extra threads if necessary.
  - Executed with extra threads if the # of key-value pairs  $\geq$  this threshold.
    - `Long.MAX_VALUE`: Suppress concurrency/parallelism. Use a single thread.
    - 1: Maximize concurrency/parallelism.
- The number of threads to be used:
  - If (# of key-value pairs) < threshold
    - No extra threads to be used
  - If (# of KV pairs) / threshold  $\geq$  (# of CPU cores)
    - Use (# of CPU cores)
    - C.f. `Runtime.availableProcessors()`
  - If (# of KV pairs) / threshold < (# of CPU cores)
    - Use (# of KV pairs) / threshold
      - Rounded to an int number

## search Bulk Operations


- Searches a key-value pair that satisfies a given **search criterion** (specified as a LE)
  - Returns a non-null result if found. Returns null otherwise.
  - Skips further search when the first search hit is returned.
- `Path path = map.search(500, (k,v)->{v>10000? v: null} )`
  - Second param: search function (`BiFunction`)
    - Is there at least one file that has more than 10,000 access count?
  - If there exists such a file, `path` contains a path to that file.
  - Only the first search hit is returned.
- `Integer count = map.searchValues(500, (v)->{v>10000? v: null} )`
  - Second param: search function (`Function`)
  - If there exists such a file, `count` contains the access count of that file.

# reduce Bulk Operations

- Reduces key-value pairs to a single value.

```
- int maxCount = map.reduceToInt(500,  
                                (k,v)-> v,  
                                0,  
                                (max, v)-> v>max? v: max );
```

- 2nd param: transformer (`BiFunction`)
  - Expected to work like a map operation
- 4th param: reducer (`BinaryOperator`)
- 3rd param: the initial value for accumulated value (`max`).



```
- int maxCount = map.reduceToInt(50,  
                                (k,v)->v,  
                                0,  
                                (max, v)-> v>max? v: max );
```

- 4th param: reducer (`BinaryOperator`)
  - 3rd param: the initial value for accumulated value (`max`).
- Generalized form of reduce operations with the Streams API

```
- T result = aStream.reduce(initValue, (result, element)-> ... );  
  
- T result = initValue;  
  for(T element: collection){  
    result = accumulate(result, element);  }
```

## Notes on Bulk Operations

- Bulk operations
  - are **thread-safe** in that they never corrupt key-value pairs.
  - may not be performed on the most up-to-date key-value pairs (weak consistency)
    - Write threads can modify key-value pairs when read threads are reading key-value pairs.
      - It is guaranteed that readers and writers do not corrupt key-value pairs.
      - c.f. Concurrent iteration

## Recap: Collection Processing with LEs


- Java 8 made major improvements to the Collection API by
  - Adding new **static and default methods** in existing interfaces
    - e.g., `Iterable.forEach()`

## Recap: External v.s. Internal Iteration

- **External** iteration: Iterate over a collection and performs an operation on each element in turn.

```
- Iterator<ArrayList> iterator = strList.iterator();
  while( iterator.hasNext() ) {
    System.out.print(iterator.next()); }
```

- Iteration occurs **outside** of a collection.
- Need to write a **boilerplate code** whenever you need to iterate over the collection.



61

62

- The loop mixes up *what you want to do on a collection* and *how you do it*.
  - May not be that maintainable because “what” is often obscured by “how.” (“How” is often emphasized too much than “what.”)


```
- Iterator<ArrayList> iterator = strList.iterator();
  while( iterator.hasNext() ) {
    System.out.print(iterator.next()); }
```

- Inherently serial
  - Hard to make it concurrent/parallel.

- **Internal** iteration:

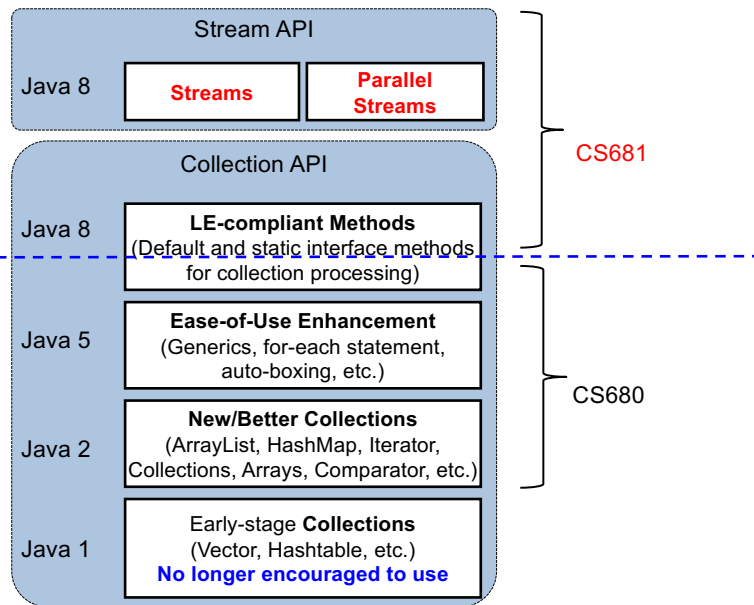
```
- int maxCount = map.reduceToInt(50,
                                (k,v)->v,
                                0,
                                (max, v)-> v>max? v: max );
```

- Does not return an `Iterator` that externally controls the iteration
- Creates an equivalent object, which works **inside** of a collection.
  - A collection internally uses the iterator-like object to perform iteration



63

64






- Client code simply states “**what**” you want to do on a collection. “**How**” is hidden.
  - Collection processing looks more **declarative**, not procedural.
    - c.f. SQL statements

- `ConcurrentHashMap` **internally** uses threads to perform bulk operations.

```
- int maxCount = map.reduceToInt(2,
                                (k,v) -> v,
                                0,
                                (max, v) -> v > max ? v : max);
```

- Client code does not have to **externally** use threads.



65

## Benefits of Using Lambda Expressions

- Can make your code more concise (less repetitive)
- Can enjoy the power of functional programming
  - e.g., higher-order functions
- Can gain a new way to access collections
  - “Internal” iteration as opposed to traditional “external” iteration
    - e.g., `Map-Reduce` data processing
- Can simplify **concurrent programming (multi-threading)** in Java

67

## Recap: Concurrent Collections

- “Ready-made” **thread-safe** collections
  - Introduced in Java 5 (2004) and enhanced in subsequent versions
    - **Queue**
      - `ConcurrentLinkedQueue` (since Java 5)
      - `ConcurrentLinkedDeque` (since Java 7)
      - `ArrayBlockingQueue` (since Java 5)
      - `LinkedBlockingQueue` (since Java 5)
      - `DelayQueue` (since Java 5)
      - `PriorirtyBlockingQueue` (since Java 5)
      - `LinkedBlockingDeque` (since Java 6)
      - `LinkedTransferQueue` (since Java 7)
    - **Map**
      - `ConcurrentHashMap` (since Java 5)
      - `ConcurrentSkipListMap` (since Java 6)
    - **Set**
      - `ConcurrentSkipListSet` (since Java 6)
  - `java.util.concurrent.CopyOnWriteXyz` classes
    - `CopyOnWriteArrayList`
    - `CopyOnWriteArraySet`

## Every Concurrent Collection...

- Implements its **public methods** in **thread-safe** and **performance-aware** manners, just like `ConcurrentHashMap`.
  - c.f. lock stripping in `ConcurrentHashMap`
- Supports **thread-safe**, **weakly-consistent iteration** with an iterator and a spliterator(s)
  - just like `ConcurrentHashMap`
- Provides **thread-safe** methods for the other types of **compound operations**
  - i.e. Navigation and check-and-act operations
  - C.f. `putIfAbsent(key, value)` and `compute()` in `ConcurrentHashMap`
- Provides **thread-safe** methods for **bulk operations**.

## Concurrent Hash Set???


- No concurrent collection class for **hash-based sets** in Java API
  - No such class like `ConcurrentHashSet`.
- You can “**emulate**” it through `ConcurrentHashMap`:
  - `Set<String> set = ConcurrentHashMap<String, Integer>.newKeySet();`
    - Values (`Integer`s) are ignored.
    - `set` is a **hash-based set** that contains a series of `String` data.
    - `newKeySet()` is a static factory method to generate a `ConcurrentHashMap.KeySetView<K, Boolean>`
      - `ConcurrentHashMap.KeySetView<String, Boolean>` in the above example.

## Some Major Concurrent Collections

- **ConcurrentLinkedQueue**
  - Concurrent implementation of `Queue`
    - FIFO (First-In-First-Out) queue
- **ConcurrentLinkedDeque**
  - Concurrent implementation of `Deque`
    - LIFO (Last-In-First-Out) queue
- **ConcurrentSkipListMap**
  - An implementation of `ConcurrentNavigableMap`.
  - Map entries are kept sorted according to the natural ordering of their keys or by a custom `Comparator`.
- **ConcurrentSkipListSet**
  - A concurrent implementation of `NavigableSet`.
  - Set elements are kept sorted according to the natural ordering or by a custom `Comparator`.

## Concurrent Array List???

- No concurrent collection class for **lists** in Java API
  - No such class like `ConcurrentArrayList`.
  - Why? Performance-wise, lists are not that great collections for multi-threaded programs
    - **Insertion and removal** are expensive for **non-tail elements**
      - c.f. Slides on List v.s. Queue and ArrayList v.s. LinkedList
- Hard to implement **index-based random access** in a concurrent and efficient manner.
  - e.g., Weakly consistent iteration cannot be more efficient than perfectly consistent iteration, if insertion/removal often occurs for non-tail elements.




# Exercise: Online Shopping Cart

– If you need/want to use an `ArrayList`, you have to craft thread-safe client code for it yourself with client-side locking.

– Try a concurrent queue class (e.g. `ConcurrentLinkedQueue`) or set class (e.g. `ConcurrentHashMap.KeySetView<K, Boolean>`) if you like to take advantages of concurrent collections.


– FYI: `CopyOnWriteArrayList`

- Concurrent collection class for `ArrayList`
- Thread-safe, but not that efficient in general (to be explained)
  - Can be efficient only in limited use cases



• Thanks to `ConcurrentLinkedQueue`, all methods of `cart` do not need thread synchronization to access `inCartItems`.

```
getSubTotal() {
    float subtotal;
    for(Product item: inCartItems){
        subtotal += item.getPrice();
    }
    return subtotal; }
// No thread synch!
// Weakly consistent iteration
```




• Can use `ConcurrentLinkedQueue<Product>` for `inCartItems`

- It is originally typed with `LinkedList`

• Thanks to `ConcurrentLinkedQueue`, all methods of `cart` do not need thread synchronization to access `inCartItems`.

```
addItem(Product item){
    inCartItems.offer(item); }
// No thread synch!
```

```
getInCartItems() {
    return inCartItems; }
// No thread synch!
```



• `removeItem()` used an index number for a `LinkedList`.

- `public void removeItem(int productIndex)`

– Need to change the method signature

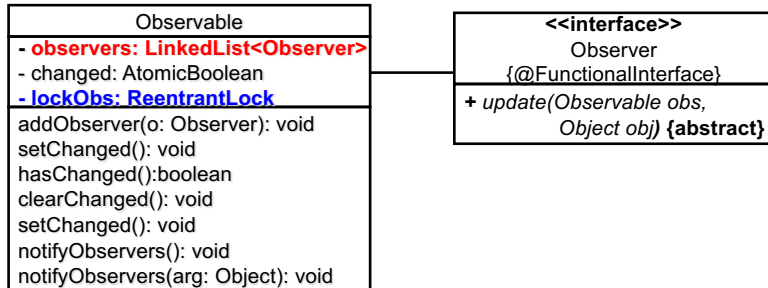
```
public void removeItem(Product item)
public void removeItem(String productId)
```

```
removeItem(Product item){
    inCartItems.remove(item); }
// No thread synch!

removeItem(String productId){
    inCartItems.removeIf(
        (Product p)->
        p.getID().equals(productId)) }
// No thread synch!
```

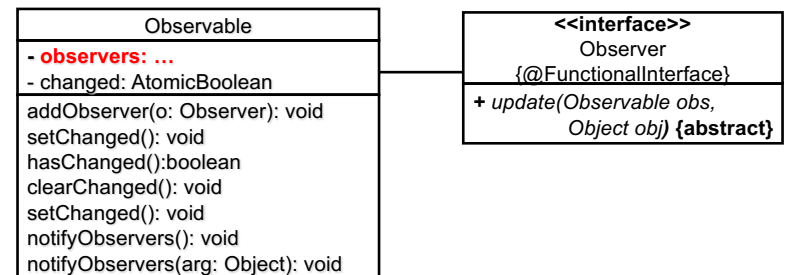
## Recap: Concurrent Observer

- C.f. HW 16
  - `lockObs` is used to guard `observers` (a `LinkedList`).



## HW 18

- Revise your HW 16 code to replace `LinkedList` with
  - a concurrent queue, or
    - `ConcurrentLinkedQueue`
  - a concurrent hash set
    - `ConcurrentHashMap.KeySetView<K, Boolean>`
- Make sure that `observable` no longer needs a `ReentrantLock` data field to guard `observers`.



## Concurrent Collections


- “Ready-made” **thread-safe** collections
  - Introduced in Java 5 (2004) and enhanced in subsequent versions
    - **Queue**
      - `ConcurrentLinkedQueue` (since Java 5)
      - `ConcurrentLinkedDeque` (since Java 7)
      - `ArrayBlockingQueue` (since Java 5)
      - `LinkedBlockingQueue` (since Java 5)
      - `DelayQueue` (since Java 5)
      - `PriorirtyBlockingQueue` (since Java 5)
      - `LinkedBlockingDeque` (since Java 6)
      - `LinkedTransferQueue` (since Java 7)
    - **Map**
      - `ConcurrentHashMap` (since Java 5)
      - `ConcurrentSkipListMap` (since Java 6)
    - **Set**
      - `ConcurrentSkipListSet` (since Java 6)
  - `java.util.concurrent.CopyOnWriteXyz` classes
    - `CopyOnWriteArrayList`
    - `CopyOnWriteArraySet`

## Appendix: Copy-On-Write (COW) Collections

# Copy-On-Write (COW) Collections

- `CopyOnWriteArrayList`
- `CopyOnWriteArraySet`
  - Concurrent replacements of synchronized wrappers for `ArrayList` and `ArraySet`.
- Key features
  - Thread-safe public methods
    - No client-side locking is necessary to call them.
  - Read threads are never mutually excluded with each other.
  - Read threads and write threads are never mutually excluded.
  - Write threads are never mutually-excluded.
- Pros
  - **No client-side locking** is necessary for iteration.
    - Read threads ARE NOT mutually excluded **with each other** and **with writer threads**.
      - Each iterator has a thread-specific snapshot of List elements.
      - Different readers get different snapshots and access them concurrently.

```
Iterator it = aCOWArrayList.iterator(); // A snapshot is created.
while( it.hasNext() )
    doSomething( it.next() );
```




Reading data

# Snapshot-based Iteration in COW Collections


- Support thread-safe, **snapshot-based iteration**.
  - A read thread references and operates on a **collection snapshot**, which is a collection that is up-to-date when an iterator is created.

```
Iterator it = aCOWArrayList.iterator(); // A snapshot is created.
while( it.hasNext() )
    doSomething( it.next() );
```




Reading data

- Pros
  - **No client-side locking** is necessary for iteration.
    - Read threads ARE NOT mutually excluded **with each other** and **with writer threads**.
      - Each iterator has a thread-specific snapshot of List elements.
      - Different readers get different snapshots and access them concurrently.




- Cons
  - The snapshot can become **outdated**.
    - e.g., when a write thread adds/removes collection elements after a snapshot is created.
  - **No consistency** preserved; Each iterator will **NOT** reflect additions, removals and other changes to the list after the iterator was created.
    - c.f. weak consistency in `ConcurrentHashMap` (and other non-COW concurrent collections)
- Trades consistency for performance

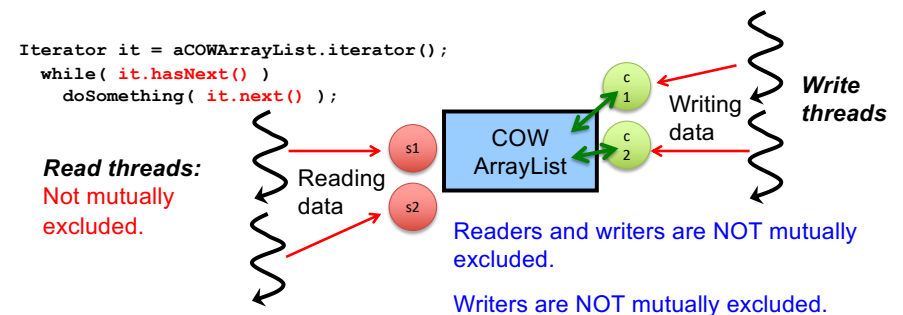
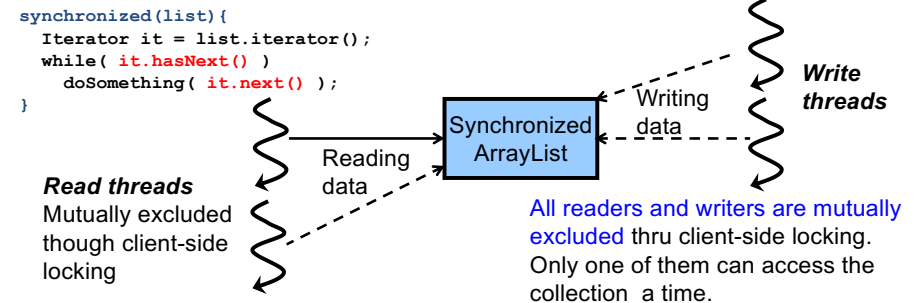




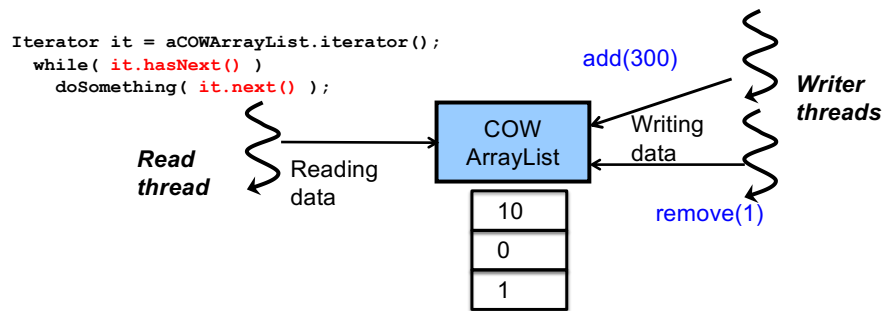
## Copy-On-Write (COW)

- Making **a copy of a collection** when a write thread updates the collection's elements
- A write thread
  - Performs `add()`, `remove()`, `set()` and other state-changing (or mutative) methods on **a duplicate copy** of collection elements.
  - Synchronizes the updated/modified copy with the original element set.
- Write threads ARE NOT mutually excluded **with each other** and **with read threads**.



- Cons
  - **No consistency** preserved. **Why?**
  - State updates (particularly, element insertion and removal) are often expensive for lists.
    - Due to index-based element ordering





- Guaranteed that collection elements are never corrupted.
  - Read thread obtains:
    - (10, 0, 1),
    - (10, 0, 1, 300)
    - (10, 0), or
    - (10, 0, 300)
  - It never obtains corrupted list such as (10, 300).

## Pros and Cons in COW Performance

- Pros
  - No concerns about data corruption.
  - Improved performance for iteration
- Cons
  - State-changing methods (e.g., add(), remove()) are very slow.
  - Never use COW collections in single-threaded programs.
  - Their overhead grows exponentially as the number of elements increases.

– The overhead of add() [msec]

| » # of elems | ArrayList | SyncArrayList | COWArrayList |
|--------------|-----------|---------------|--------------|
| » 1,000      | 0         | 0             | 14           |
| » 5,000      | 0         | 0             | 102          |
| » 10,000     | 0         | 0             | 409          |
| » 20,000     | 0         | 0             | 1,712        |
| » 30,000     | 15        | 16            | 4,566        |

## When to Use COW Collections?

- Read operations are executed a lot more often than write operations.
- When the # of read threads is a lot greater than the # of write threads.
- When state-changing methods are rarely called.
- When the # of elements is relatively small.

## In Summary...

- Be aware of the characteristics of COW collections.
- Be conservative to use them.